

Loss modelling and reserving analytics in R

A hands-on workshop

Katrien Antonio & Jonas Crevecoeur

Arcturus 2020 workshop | November 25-27, 2020

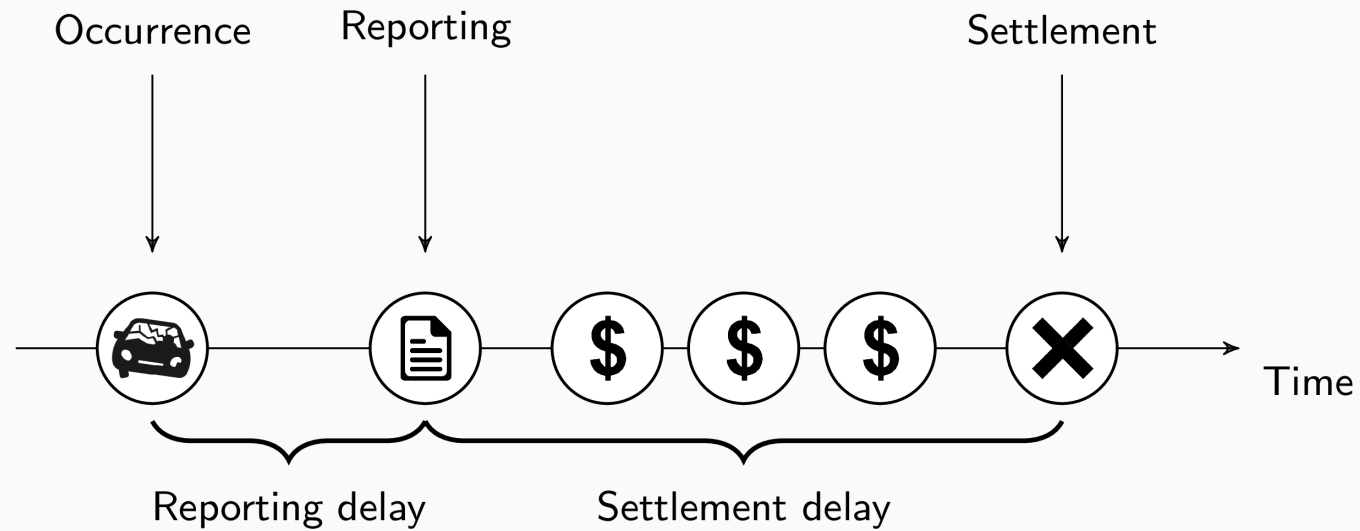
Today's Outline

- Motivation and strategies
- Reserving data structures (part 1)
 - IBNR and RBNS reserve
 - daily and yearly reserving data
- Reserving data structures (part 2)
 - individual and aggregated reserving data
 - runoff triangles
- Claims reserving with triangle
 - chainladder model
 - {ChainLadder} package
 - chainLadder GLM implementation
- When the chain ladder method fails
 - detection tools for triangle stability
 - creating homogeneous triangles
- Research outlook
 - granular reserving methods for predicting the IBNR reserve
 - individual reserving methods for predicting the RBNS reserve

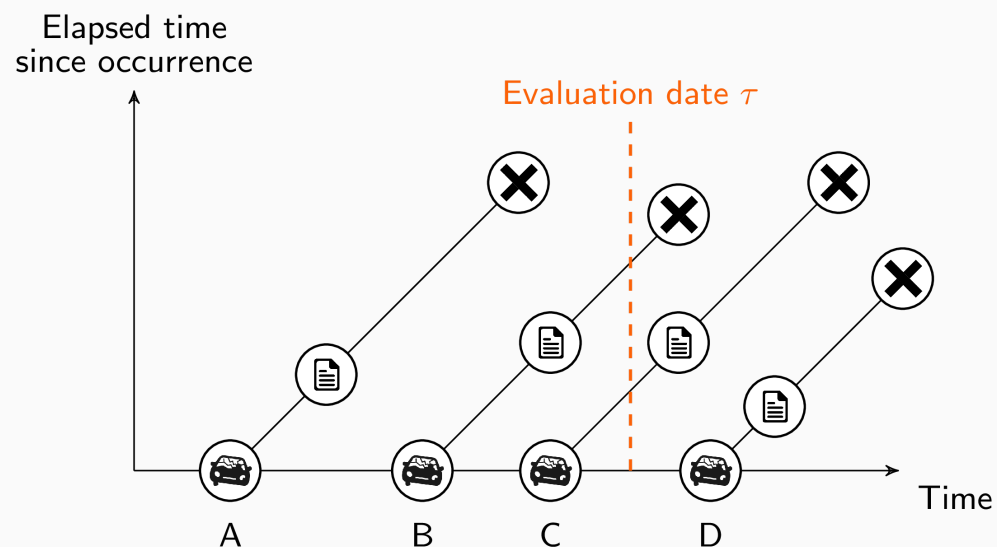
Motivation and strategies

Motivation

Insurers observe the **detailed evolution** of **individual claims**.



Motivation (continued)



Observed claims are **censored** due to **delays** (reporting, settlement) in the claim development process

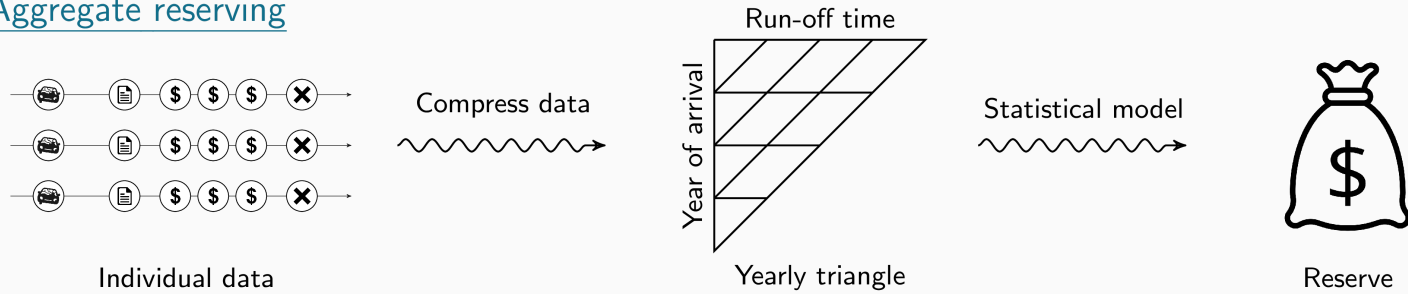
Reserve: future costs for claims that occurred in **past exposure periods** (claim B and C)

- RBNS reserve: Reserve for Reported, But Not yet Settled claims (claim B)
- IBNR reserve: Reserve for Incurred, But Not yet Reported claims (claim C)

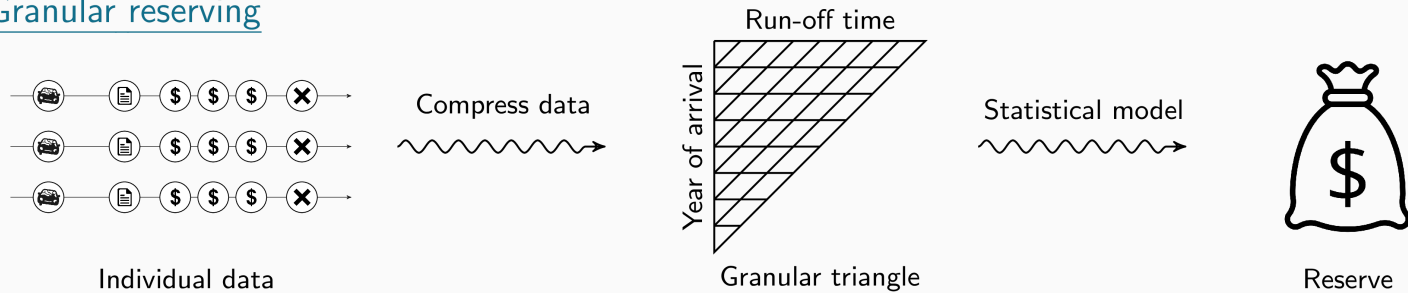
Pricing: all costs for claims that will occur in **future** insured **exposure periods** (claim D)

Three strategies for non-life reserving

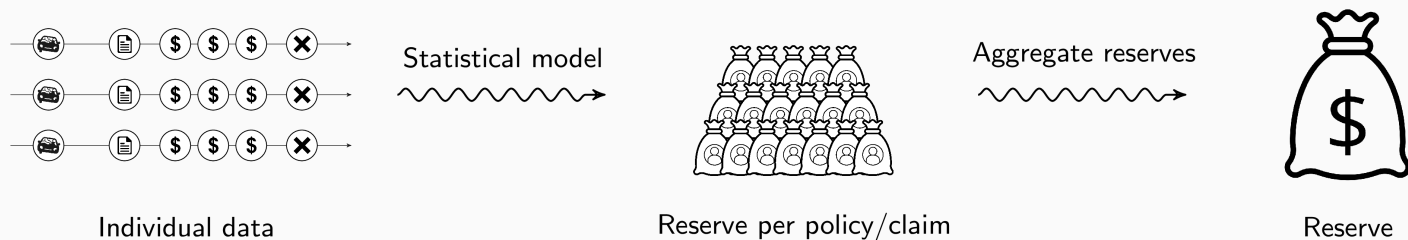
Aggregate reserving



Granular reserving



Individual reserving



Three strategies for non-life reserving (continued)

Aggregate reserving

First aggregate the past claim history into a **small number of summary statistics**, then predict the total reserve based on these summary statistics.

Granular reserving

First aggregate the past claim history into a **large number of summary statistics**, then predict the total reserve based on these summary statistics.

Individual reserving

First predict the future costs for **individual claims**, then aggregate these individual predictions to predict the total reserve.

Data set used in this session

We illustrate reserving analytics based on a **simulated dataset** registering the **detailed development** of 30,000 claims between January 1, 2010 and December 31, 2019.

- Information is available at daily level
- Records correspond to payments made for these claims

The data is stored in a .RData file in the `data` directory in the course material:

```
# install.packages("rstudioapi")
dir <- dirname(rstudioapi::getActiveDocumentContext())$path
setwd(dir)

load("data/reserving_data.RData")
```

##	accident_number	accident_date	reporting_date	settlement_date	reporting_delay	settlement_delay	payment_date	payment_size
## 1	1	2012-08-27	2012-09-01	2012-12-03	5	98	2012-09-01	0.0000
## 2	1	2012-08-27	2012-09-01	2012-12-03	5	98	2012-09-07	446.2251
## 3	1	2012-08-27	2012-09-01	2012-12-03	5	98	2012-09-18	490.1792
## 4	1	2012-08-27	2012-09-01	2012-12-03	5	98	2012-10-01	587.7409
## 5	1	2012-08-27	2012-09-01	2012-12-03	5	98	2012-10-19	1726.0902
## 6	2	2013-09-20	2013-09-26	2014-04-06	6	198	2013-09-26	0.0000

Data set used in this session (continued)

Discretize daily dates to yearly indices:

```
date_to_year <- function(date, base_year) {  
  year <- as.numeric(format(date, '%Y')) - base_year  
}  
  
reserving_data <- reserving_data %>%  
  mutate(accident_year = date_to_year(accident_date, 2010),  
         reporting_year = date_to_year(reporting_date, 2010),  
         payment_year = date_to_year(payment_date, 2010),  
         development_year = payment_year - accident_year + 1,  
         settlement_year = date_to_year(settlement_date, 2010))
```

```
##  accident_number accident_date reporting_date settlement_date reporting_delay settlement_delay payment_date payment_size accident_year reporting_year payment_year  
## 1              1  2012-08-27  2012-09-01  2012-12-03             5              98  2012-09-01      0.0000             2              2              2  
## 2              1  2012-08-27  2012-09-01  2012-12-03             5              98  2012-09-07    446.2251             2              2              2  
## 3              1  2012-08-27  2012-09-01  2012-12-03             5              98  2012-09-18    490.1792             2              2              2  
## 4              1  2012-08-27  2012-09-01  2012-12-03             5              98  2012-10-01    587.7409             2              2              2  
## 5              1  2012-08-27  2012-09-01  2012-12-03             5              98  2012-10-19   1726.0902             2              2              2  
## 6              2  2013-09-20  2013-09-26  2014-04-06             6             198  2013-09-26      0.0000             3              3              3  
##  development_year settlement_year  
## 1              1              2  
## 2              1              2  
## 3              1              2  
## 4              1              2  
## 5              1              2  
## 6              1              4
```



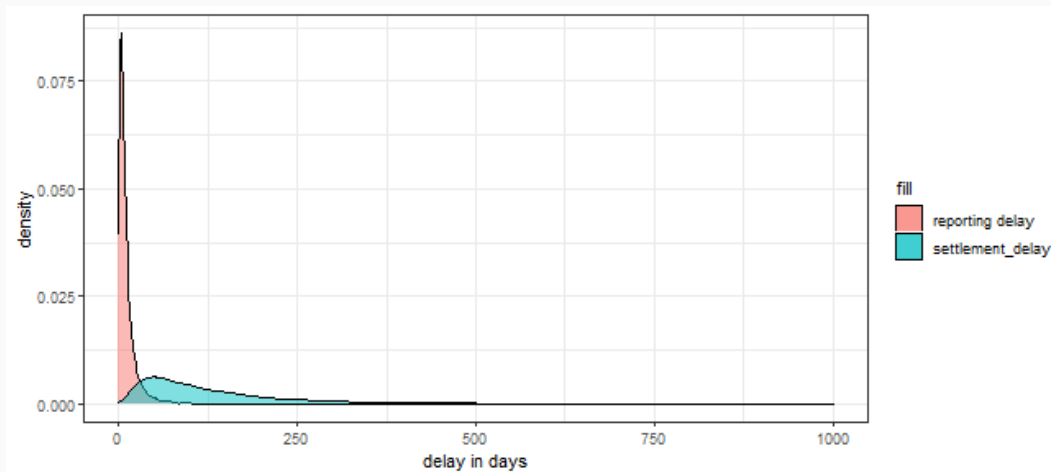
Your turn

In this warm up exercise you load the `reserving data` data set and get some feel for the data.

1. Visualize the reporting and settlement delay of claims with a density plot.
2. When was the last payment registered in the data set?
What do you conclude?
3. What is the average number of payments per claim?
4. Calculate the number of claims per accident year.

For **Q.1** we plot the density of reporting and settlement delay:

```
ggplot(data = reserving_data) +  
  theme_bw() +  
  geom_density(aes(reporting_delay, fill = 'reporting delay'),  
               alpha = .5) +  
  geom_density(aes(settlement_delay, fill = 'settlement_delay'),  
               alpha = .5) +  
  xlab('delay in days') +  
  xlim(c(0, 1000))
```



For **Q.2** the last payment was in 2024, i.e. the simulated data is not yet censored!

```
max(reserving_data$payment_date)
```

```
## [1] "2024-05-05"
```

For **Q.3** the data set records one row per payment.

```
num_claim <- length(unique(reserving_data$accident_number))
num_payment <- nrow(reserving_data)

num_payment/num_claim
```

```
## [1] 3.425433
```

At reporting, the data set registers a zero payment for each claim. The average number of non-zero payments per claim is

```
num_payment <- sum(reserving_data$payment_size > 0)
num_payment/num_claim
```

```
## [1] 2.376633
```

For **Q.4** we group the data by accident year and summarise the number of claims.

```
reserving_data %>%
  group_by(accident_year) %>%
  summarise(num_claims = n())
```

```
## # A tibble: 10 x 2
##   accident_year num_claims
##           <dbl>     <int>
## 1             0      10275
## 2             1      10108
## 3             2      10633
## 4             3      10295
## 5             4       10213
## 6             5      10130
## 7             6      10222
## 8             7       10069
## 9             8      10261
## 10            9      10557
```

The portfolio is stable over time.

Reserving data structures: daily and yearly data

From daily to yearly data

Reserves are often calculated based on yearly, **aggregated data**.

Data aggregation removes the daily noise from the data.

In a first step, we construct a data set that consists of **one record per claim and per development year**.

We construct this data set by merging three partial data sets:

- `records`: **empty data set** with one record per claim per development year
- `claim_data`: data set with one record per claim containing the **static claim covariates** (e.g. accident year)
- `payment_data`: data set containing the total **amount paid per claim and per development year**.

This data set contains only records for development years with at least one payment.

From daily to yearly data (continued)

For `records`, `expand.grid` creates a data.frame with all combinations of `accident_number` and `development_year`.

```
max_dev_year <- 10
accidents <- unique(reserving_data$accident_number)

records <- expand.grid(accident_number = accidents,
                      development_year = 1:max_dev_year)
```

```
head(records)
```

```
##   accident_number development_year
## 1             1             1
## 2             2             1
## 3             3             1
## 4             4             1
## 5             5             1
## 6             6             1
```

`expand.grid` creates a data.frame with all possible combinations of the input vectors.

```
expand.grid(A = c(1, 5, 7), B = c(4, 6))
```

```
##   A B
## 1 1 4
## 2 5 4
## 3 7 4
## 4 1 6
## 5 5 6
## 6 7 6
```

From daily to yearly data (continued)

For `claim_data`, we select for each `accident_number` the first row in `reserving_data`.

```
claim_data <- reserving_data %>%
  group_by(accident_number) %>%
  slice(1) %>%
  ungroup() %>%
  select(accident_number, accident_year,
         accident_date, reporting_year,
         settlement_year)
```

```
head(claim_data)
```

```
## # A tibble: 6 x 5
##   accident_number accident_year accident_date reporting_year
##           <int>         <dbl> <date>             <dbl>
## 1             1             2 2012-08-27             2
## 2             2             3 2013-09-20             3
## 3             3             5 2015-09-24             5
## 4             4             9 2019-01-30             9
## 5             5             2 2012-01-07             2
## 6             6             8 2018-12-25             8
```

`slice(k)` selects the k -th record of a data set

```
df <- data.frame(x = c(1, 2, 3, 4, 5, 6), y = c(1, 1, 1, 2, 2,
df %>% slice(4)
```

```
##   x y
## 1 4 2
```

In combination with `group_by` we can select the k -th record within each group

```
df %>% group_by(y) %>% slice(1)
```

```
## # A tibble: 2 x 2
## # Groups:   y [2]
##       x     y
##   <dbl> <dbl>
## 1     1     1
## 2     4     2
```


From daily to yearly data (continued)

For `payment_data`, we aggregate payments per claim by `development_year`.

```
payment_data <- reserving_data %>%  
  group_by(claim_number, development_year) %>%  
  summarise(size = sum(payment_size),  
            payment = size > 0)
```

```
head(payment_data)
```

```
## # A tibble: 6 x 4  
## # Groups:   claim_number [5]  
##   claim_number development_year size payment  
##           <int>           <dbl> <dbl> <lgl>  
## 1             1             1 3250. TRUE  
## 2             2             1 2347. TRUE  
## 3             2             2 1189. TRUE  
## 4             3             1 1975. TRUE  
## 5             4             1 3107. TRUE  
## 6             5             1 2376. TRUE
```

`summarise` returns one record for each combination of `claim_number` and `development_year` in `reserving_data`.

The data set contains no records development years in which no payment was made for a claim!

From daily to yearly data (continued)

Merge `records`, `claim_data` and `payment_data` into a new data set `individual_data`.

```
individual_data <- records %>%  
  left_join(claim_data,  
            by = 'accident_number') %>%  
  left_join(payment_data,  
            by = c('accident_number', 'development_year')) %>%  
  mutate(size = replace_na(size, 0),  
         payment = replace_na(payment, FALSE))
```

```
dim(individual_data)
```

```
## [1] 300000      8
```

```
head(individual_data, 3)
```

##	accident_number	development_year	accident_year	accident_date	reporting_year	settlement_year	size	payment
## 1	1	1	2	2012-08-27	2	2	3250.235	TRUE
## 2	2	1	3	2013-09-20	3	4	2347.445	TRUE
## 3	3	1	5	2015-09-24	5	5	1975.269	TRUE

`left_join` combines records in the first (left) with records from the second (right) data set with matching values in the `by` column.

If no match is found in the right data set an `NA` value is imputed.

```
left <- data.frame(x = c(1, 2, 3), left = 1)  
right <- data.frame(x = c(1, 2, 4), right = 1)  
left_join(left, right, by = 'x')
```

##	x	left	right
## 1	1	1	1
## 2	2	1	1
## 3	3	1	NA

Censoring the data

The simulated data is not censored, i.e. we observe the full development of all claims that occur before 2019. In practice the available data set is censored and we only observe:

- reported claims
- development information for past calendar years

We censor the data assuming that we observe data until the end of 2019

```
observed_data <- individual_data %>%  
  mutate(calendar_year = accident_year + development_year - 1) %>%  
  filter(calendar_year >= reporting_year,  
         calendar_year <= 9)  
  
unobserved_data <- individual_data %>%  
  mutate(calendar_year = accident_year + development_year - 1) %>%  
  filter(calendar_year > 9)
```

IBNR and RBNS reserves

Reserving models predict the total payments in the unobserved data

```
reserve_actual <- sum(unobserved_data$size)
reserve_actual
```

```
## [1] 2468246
```

Some reserving methods split this reserve in an

- IBNR reserve: a reserve for Incurred, But Not (yet) Reported claims
- RBNS reserve: a reserve for Reported, But Not (yet) Settled claims

```
unobserved_data %>%
  mutate(reported = (reporting_year <= 9)) %>%
  group_by(reported) %>%
  summarise(reserve = sum(size))
```

```
## # A tibble: 2 x 2
##   reported reserve
##   <lgl>      <dbl>
## 1 FALSE    158665.
## 2 TRUE     2309581.
```

Reserving data structures: individual and aggregated data

From individual to aggregated data

Individual triangle

The data set `individual_data` describes the development of **individual claims discretised by calendar year**.

We can represent the claim characteristics in this data set (settlement, payment, payment size) with **two-dimensional tables** in which the rows represent **individual claims** and columns represent **development years**.

For payment size, this table consists of cells $U_{k,j}^{\text{size}}$ with $U_{k,j}^{\text{size}}$ denoting the total amount paid for claim k in development year j .

Runoff triangle

Many reserving models go one step further and remove individual claim characteristics by aggregating claims by occurrence year. The data is then represented in a (small) two-dimensional table, the so-called **incremental runoff triangle**.

$$X_{i,j}^{\text{size}} = \sum_{\text{occ.year}(k)=i} U_{k,j}^{\text{size}}.$$

As a result of the Law of Large Numbers (LLN), aggregating data from many claims into runoff triangles **reduces the variance** when all claims are independent realizations from the same distribution.

From individual to aggregated data (continued)

We construct a runoff triangle by aggregating individual claims by `accident_year`.

```
observed_data %>%
  group_by(accident_year, development_year) %>%
  summarise(value = sum(size)) %>%
  pivot_wider(values_from = value,
              names_from = development_year,
              names_prefix = 'DY.')
```

```
## # A tibble: 10 x 11
## # Groups:   accident_year [10]
##   accident_year  DY.1    DY.2    DY.3    DY.4    DY.5    DY.6    DY.7    DY.8    DY.9    DY.10
##   <dbl>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1         0 4532915. 1739699. 121876. 25952.  9305.  6703.  9479.  6217.  3468.    NA
## 2         1 4414954. 1673577. 107415. 19472.    0      0      0      0      0      NA
## 3         2 4786833. 1785616. 149240. 59348. 14858. 4404.    0      0     NA    NA
## 4         3 4653451. 1614053. 127900. 36214.  9301.    0      0     NA    NA    NA
## 5         4 4321019. 1909085. 159247. 49184. 11728.    0     NA    NA    NA    NA
## 6         5 4417129. 1722151. 103907. 65592. 15129.    NA    NA    NA    NA    NA
## 7         6 4538329. 1718509. 146138. 19248.    NA    NA    NA    NA    NA    NA
## 8         7 4578285. 1618201.  77054.    NA    NA    NA    NA    NA    NA    NA
## 9         8 4503558. 1756243.    NA    NA    NA    NA    NA    NA    NA    NA
## 10        9 4575529.    NA    NA    NA    NA    NA    NA    NA    NA    NA
```

`pivot_wider` widens the data set by increasing the number of columns, while decreasing the number of rows.

`names_from`: The resulting data set contains a separate column for each unique outcome of this variable.

`values_from`: Values from this variable are stored in the newly created columns.

```
df <- data.frame(x = c(1, 1, 1, 2, 2, 2, 3, 3, 3),
                 y = c(1, 2, 3, 1, 2, 3, 1, 2, 3),
                 value = c(1, 2, 3, 4, 5, 6, 7, 8, 9))

pivot_wider(df,
             names_from = y,
             values_from = value)
```

```
## # A tibble: 3 x 4
##       x     1     2     3
##   <dbl> <dbl> <dbl> <dbl>
## 1     1     1     2     3
## 2     2     4     5     6
## 3     3     7     8     9
```

From individual to aggregated data (continued)

A more sophisticated function for computing incremental runoff triangles is available in the R code.

```
incremental_triangle(observed_data,  
                     variable = 'payment',  
                     lower_na = TRUE)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]  
## [1,] 2264 550  20   5   2   1   1   1   1   0  
## [2,] 2269 533  27   5   0   0   0   0   0  NA  
## [3,] 2307 557  26   8   3   1   0   0  NA  NA  
## [4,] 2289 537  23   8   1   0   0  NA  NA  NA  
## [5,] 2239 581  33   8   1   0  NA  NA  NA  NA  
## [6,] 2211 532  24   6   3  NA  NA  NA  NA  NA  
## [7,] 2197 537  30   5  NA  NA  NA  NA  NA  NA  
## [8,] 2250 510  18  NA  NA  NA  NA  NA  NA  NA  
## [9,] 2271 544  NA  NA  NA  NA  NA  NA  NA  NA  
## [10,] 2319  NA  NA  NA  NA  NA  NA  NA  NA  NA
```

Using this function you can easily inspect the reserving data using multiple triangles.

```
incremental_triangle(observed_data,  
                     variable = 'size',  
                     lower_na = TRUE)
```

```
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]      [,8]      [,9]  
## [1,] 4532915 1739699 121876.31 25951.66 9304.915 6703.063 9478.77 6217.266 3468.364  
## [2,] 4414954 1673577 107415.07 19472.27  0.000  0.000  0.00  0.000  0.000  
## [3,] 4786833 1785616 149239.81 59348.18 14857.743 4404.145  0.00  0.000  NA  
## [4,] 4653451 1614053 127900.21 36213.57 9301.109  0.000  0.00  NA  NA  
## [5,] 4321019 1909085 159246.74 49184.25 11727.691  0.000  NA  NA  NA  
## [6,] 4417129 1722151 103906.89 65592.12 15129.218  NA  NA  NA  NA  
## [7,] 4538329 1718509 146138.00 19248.11  NA  NA  NA  NA  NA  
## [8,] 4578285 1618201  77053.93  NA  NA  NA  NA  NA  NA  
## [9,] 4503558 1756243  NA  NA  NA  NA  NA  NA  NA  
## [10,] 4575529  NA  NA  NA  NA  NA  NA  NA  NA
```


Cumulative runoff triangles

Cumulative runoff triangles are constructed by computing the cumulative row sum of an **incremental runoff triangle**, i.e.

$$C_{i,j}^{\text{size}} = \sum_{l=1}^j X_{i,l}^{\text{size}}.$$

When using a new reserving method, check whether the incremental or cumulative runoff triangle is required!

```
incremental_triangle(observed_data,  
                     variable = 'size',  
                     lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	4532915	1739699	121876.31	25951.66	9304.915	6703.063	9478.77	6217.266	3468.364
##	[2,]	4414954	1673577	107415.07	19472.27	0.000	0.000	0.00	0.000	0.000
##	[3,]	4786833	1785616	149239.81	59348.18	14857.743	4404.145	0.00	0.000	NA
##	[4,]	4653451	1614053	127900.21	36213.57	9301.109	0.000	0.00	NA	NA
##	[5,]	4321019	1909085	159246.74	49184.25	11727.691	0.000	NA	NA	NA
##	[6,]	4417129	1722151	103906.89	65592.12	15129.218	NA	NA	NA	NA
##	[7,]	4538329	1718509	146138.00	19248.11	NA	NA	NA	NA	NA
##	[8,]	4578285	1618201	77053.93	NA	NA	NA	NA	NA	NA
##	[9,]	4503558	1756243	NA	NA	NA	NA	NA	NA	NA
##	[10,]	4575529	NA	NA	NA	NA	NA	NA	NA	NA

```
cumulative_triangle(observed_data,  
                    variable = 'size',  
                    lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	4532915	6272614	6394490	6420442	6429747	6436450	6445929	6452146	6455614	6455614
##	[2,]	4414954	6088531	6195946	6215419	6215419	6215419	6215419	6215419	6215419	NA
##	[3,]	4786833	6572449	6721689	6781037	6795895	6800299	6800299	6800299	NA	NA
##	[4,]	4653451	6267504	6395404	6431618	6440919	6440919	6440919	NA	NA	NA
##	[5,]	4321019	6230104	6389351	6438535	6450263	6450263	NA	NA	NA	NA
##	[6,]	4417129	6139280	6243187	6308779	6323908	NA	NA	NA	NA	NA
##	[7,]	4538329	6256837	6402975	6422224	NA	NA	NA	NA	NA	NA
##	[8,]	4578285	6196486	6273540	NA	NA	NA	NA	NA	NA	NA
##	[9,]	4503558	6259801	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	4575529	NA	NA	NA	NA	NA	NA	NA	NA	NA

Claims reserving with triangles

Mack chain ladder method

Assumptions Mack (1993) **chain ladder method**:

1. There exist development factors f_j such that

$$E(C_{i,j+1} \mid C_{i,1}, \dots, C_{i,j}) = C_{i,j} \cdot f_j \quad \text{for } 1 \leq i \leq l, 1 \leq j \leq l-1,$$

with l the dimension of the runoff triangle.

2. There exist parameters σ_j such that

$$\text{Var}(C_{i,j+1} \mid C_{i,1}, \dots, C_{i,j}) = C_{i,j} \cdot \sigma_j^2 \quad \text{for } 1 \leq i \leq l, 1 \leq j \leq l-1.$$

3. Occurrence years are independent.

Remarks:

- Method based on the **cumulative triangle**!
- Assumption 1. and 2. add an **Markov property**. Future evolutions of the reserve depend only on the last known situation.

Mack chain ladder method (continued)

Estimating the **development factors** f_j^{\wedge} :

$$f_j^{\wedge} = \frac{\sum_{i=1}^{l-j} C_{i,j+1}}{\sum_{i=1}^{l-j} C_{i,j}}.$$

```
triangle <- cumulative_triangle(observed_data, variable = 'size')
l <- nrow(triangle)
f <- rep(0, l-1)

for(j in 1:(l-1)) {
  f[j] <- sum(triangle[1:(l-j), j+1]) / sum(triangle[1:(l-j), j])
}

f

## [1] 1.381312 1.019846 1.006146 1.001563 1.000344 1.000366 1.000319 1.000274 1.000000
```

Development factors should converge to 1 for high development years.

Mack chain ladder method (continued)

Use the development factors to estimate cells in the lower triangle ($i+j > l+1$):

$$\hat{C}_{ij} = \hat{C}_{i,j-1} \cdot \hat{f}_{j-1} \quad \text{for } i+j > l+1.$$

```
triangle_completed <- triangle
for(j in 2:l) {
  triangle_completed[l:(l-j+2), j] <- triangle_completed[l:(l-j+2), j-1] * f[j-1]
}
```

Completed cumulative triangle:

triangle_completed

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	4532915	6272614	6394490	6420442	6429747	6436450	6445929	6452146	6455614	6455614
##	[2,]	4414954	6088531	6195946	6215419	6215419	6215419	6215419	6215419	6215419	6215419
##	[3,]	4786833	6572449	6721689	6781037	6795895	6800299	6800299	6800299	6802161	6802161
##	[4,]	4653451	6267504	6395404	6431618	6440919	6440919	6440919	6442977	6444741	6444741
##	[5,]	4321019	6230104	6389351	6438535	6450263	6450263	6452624	6454685	6456452	6456452
##	[6,]	4417129	6139280	6243187	6308779	6323908	6326081	6328397	6330418	6332152	6332152
##	[7,]	4538329	6256837	6402975	6422224	6432261	6434470	6436826	6438882	6440645	6440645
##	[8,]	4578285	6196486	6273540	6312100	6321965	6324137	6326452	6328473	6330206	6330206
##	[9,]	4503558	6259801	6384034	6423273	6433312	6435522	6437878	6439934	6441698	6441698
##	[10,]	4575529	6320235	6445667	6485285	6495420	6497652	6500030	6502107	6503887	6503887

Completed incremental triangle:

```
require(ChainLadder)
cum2incr(triangle_completed)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	4532915	1739699	121876.31	25951.66	9304.915	6703.063	9478.770	6217.266	3468.36
##	[2,]	4414954	1673577	107415.07	19472.27	0.000	0.000	0.000	0.000	0.00
##	[3,]	4786833	1785616	149239.81	59348.18	14857.743	4404.145	0.000	0.000	1861.91
##	[4,]	4653451	1614053	127900.21	36213.57	9301.109	0.000	0.000	2057.632	1764.07
##	[5,]	4321019	1909085	159246.74	49184.25	11727.691	0.000	2361.270	2061.371	1767.28
##	[6,]	4417129	1722151	103906.89	65592.12	15129.218	2172.474	2315.810	2021.685	1733.26
##	[7,]	4538329	1718509	146138.00	19248.11	10037.169	2209.697	2355.488	2056.324	1762.95
##	[8,]	4578285	1618201	77053.93	38559.90	9865.059	2171.807	2315.098	2021.064	1732.72
##	[9,]	4503558	1756243	124232.58	39239.04	10038.809	2210.058	2355.873	2056.660	1763.24
##	[10,]	4575529	1744706	125431.95	39617.87	10135.726	2231.395	2378.618	2076.516	1780.26

Mack chain ladder method (continued)

The estimated reserve is the sum of the estimates in the **lower half** of the **incremental runoff triangle**.

```
triangle_completed_incr <- cum2incr(triangle_completed)
lower_triangle <- row(triangle_completed_incr) + col(triangle_completed_incr) > l+1

reserve_cl <- sum(triangle_completed_incr[lower_triangle])

data.frame(reserve_cl = reserve_cl,
            reserve_actual = reserve_actual,
            difference = reserve_cl - reserve_actual,
            relative_difference_pct = (reserve_cl - reserve_actual) / reserve_actual * 100)
```

```
##   reserve_cl reserve_actual difference relative_difference_pct
## 1      2205459      2468246  -262787.3          -10.64672
```

Mack chain ladder method (continued)

These calculations are already implemented in the {ChainLadder} package.

```
require(ChainLadder)
triangle <- cumulative_triangle(observed_data, variable = 'size')
MackChainLadder(triangle)
```

```
## MackChainLadder(Triangle = triangle)
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1  6,455,614      1.000 6,455,614         0         0      NaN
## 2  6,215,419      1.000 6,215,419         0      1,243      Inf
## 3  6,800,299      1.000 6,802,161      1,862      3,366      1.8079
## 4  6,440,919      0.999 6,444,741      3,822      5,265      1.3777
## 5  6,450,263      0.999 6,456,452      6,190      7,483      1.2089
## 6  6,323,908      0.999 6,332,152      8,243      8,128      0.9860
## 7  6,422,224      0.997 6,440,645     18,422     10,090      0.5477
## 8  6,273,540      0.991 6,330,206     56,666     22,280      0.3932
## 9  6,259,801      0.972 6,441,698    181,896     35,842      0.1970
## 10 4,575,529      0.704 6,503,887    1,928,358    137,089      0.0711
##
##              Totals
## Latest:    62,217,515.46
## Dev:              0.97
## Ultimate: 64,422,974.02
## IBNR:         2,205,458.56
## Mack.S.E      148,272.46
## CV(IBNR):       0.07
```

Latest: Amount already paid

ultimate: Estimated total amount (= amount paid + reserve)

IBNR: Estimated total reserve, IBNR + RBNS!

Mack.S.E: Estimated standard deviation of the reserve

GLM chain ladder method

The chain ladder reserve estimate can also be obtained by assuming a Poisson distribution for the incremental runoff triangle with multiplicative mean, i.e.

$$X_{ij} \sim \text{POI}(\alpha_i \cdot \beta_j), \quad 1 \leq i, j \leq l.$$

```
triangle <- incremental_triangle(observed_data,  
                                variable = 'size')
```

```
triangle_long <- data.frame(  
  occ.year = as.numeric(row(triangle)),  
  dev.year = as.numeric(col(triangle)),  
  size = as.numeric(triangle))
```

```
head(triangle_long)
```

##	occ.year	dev.year	size
## 1	1	1	4532915
## 2	2	1	4414954
## 3	3	1	4786833
## 4	4	1	4653451
## 5	5	1	4321019
## 6	6	1	4417129

GLM chain ladder method (continued)

The chain ladder reserve estimate can also be obtained by assuming a Poisson distribution for the incremental runoff triangle with multiplicative mean, i.e.

$$X_{i,j} \sim \text{POI}(\alpha_i \cdot \beta_j), \quad 1 \leq i, j \leq l.$$

This Poisson model can be estimated using the `glm` (Generalized Linear Model) routine in R:

$$\log(E(X_{i,j})) = \log(\alpha_i) + \log(\beta_j).$$

```
fit <- glm(size ~ factor(occ.year) + factor(dev.year),
           data = triangle_long,
           family = poisson(link = log))
```

`factor`: Treat the input as a categorical instead of a numeric variable.

```
summary(fit)
```

```
##
## Call:
## glm(formula = size ~ factor(occ.year) + factor(dev.year), family = poisson(link = log
##      data = triangle_long)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -139.185   -64.152   -4.062    38.541   132.371
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.533e+01  4.029e-04 38048.572 < 2e-16 ***
## factor(occ.year)2 -3.792e-02  5.620e-04  -67.473 < 2e-16 ***
## factor(occ.year)3  5.229e-02  5.495e-04   95.155 < 2e-16 ***
## factor(occ.year)4 -1.686e-03  5.570e-04   -3.027 0.002472 **
## factor(occ.year)5  1.298e-04  5.568e-04    0.233 0.815596
## factor(occ.year)6 -1.931e-02  5.596e-04  -34.510 < 2e-16 ***
## factor(occ.year)7 -2.322e-03  5.574e-04   -4.165 3.12e-05 ***
## factor(occ.year)8 -1.962e-02  5.608e-04  -34.979 < 2e-16 ***
## factor(occ.year)9 -2.158e-03  5.615e-04   -3.843 0.000121 ***
## factor(occ.year)10 7.450e-03  6.171e-04   12.072 < 2e-16 ***
## factor(dev.year)2 -9.641e-01  2.982e-04 -3233.538 < 2e-16 ***
## factor(dev.year)3 -3.597e+00  1.017e-03 -3537.016 < 2e-16 ***
## factor(dev.year)4 -4.749e+00  1.915e-03 -2480.500 < 2e-16 ***
## factor(dev.year)5 -6.112e+00  4.076e-03 -1499.723 < 2e-16 ***
```

GLM chain ladder method

Predict the cells in the lower triangle

```
lower_triangle <- triangle_long$occ.year +  
  triangle_long$dev.year > 1 + 1  
  
triangle_long$size[lower_triangle] <-  
  predict(fit, newdata = triangle_long[lower_triangle, ], type  
  
triangle_long %>%  
  pivot_wider(values_from = size,  
              names_from = dev.year,  
              names_prefix = 'DY.')
```

```
## # A tibble: 10 x 11  
##   occ.year    DY.1    DY.2    DY.3    DY.4    DY.5    DY.6    DY.7    DY.8    DY.9    DY.10  
##   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1      1 4532915. 1739699. 121876. 25952.  9305.  6703.  9479.  6217.  3468.  0  
## 2      2 4414954. 1673577. 107415. 19472.    0      0      0      0      0 0.000239  
## 3      3 4786833. 1785616. 149240. 59348. 14858. 4404.    0      0 1862. 0.000261  
## 4      4 4653451. 1614053. 127900. 36214.  9301.    0      0 2058. 1764. 0.000247  
## 5      5 4321019. 1909085. 159247. 49184. 11728.    0 2361. 2061. 1767. 0.000248  
## 6      6 4417129. 1722151. 103907. 65592. 15129. 2172. 2316. 2022. 1733. 0.000243  
## 7      7 4538329. 1718509. 146138. 19248. 10037. 2210. 2355. 2056. 1763. 0.000247  
## 8      8 4578285. 1618201.  77054. 38560.  9865. 2172. 2315. 2021. 1733. 0.000243  
## 9      9 4503558. 1756243. 124233. 39239. 10039. 2210. 2356. 2057. 1763. 0.000247  
## 10     10 4575529. 1744706. 125432. 39618. 10136. 2231. 2379. 2077. 1780. 0.000250
```

Compute the reserve

```
reserve_glm <- sum(triangle_long$size[lower_triangle])  
reserve_glm
```

```
## [1] 2205459
```



Your turn

In certain lines of business, the **Corona pandemic** and social distancing measures have **reduced claim frequency**.

The R script contains code to simulate a 50% reduction in claim frequency in the months March, April and May of 2019. The new data sets are named `observed_data_covid` and `unobserved_data_covid`.

In this exercise you will investigate the effect of this reduction in claim frequency on the accuracy of the chain ladder method.

1. Compute the actual reserve from the covid data set.
2. Estimate the reserve using the chain ladder method from the {ChainLadder} package.
3. Compute the difference between the estimated and actual reserve. How many standard deviations is this error?

For **Q.1** we compute the actual reserve from the unobserved data.

```
reserve_covid_actual <- sum(unobserved_data_covid$size)
reserve_covid_actual
```

```
## [1] 2389277
```

The reserve is smaller since less claims have occurred in the past. However, the difference is small since many claims from these months have already settled.

```
c(reserve_actual, reserve_covid_actual)
```

```
## [1] 2468246 2389277
```

The effect on the amount already paid is larger.

```
c(already_paid_no_covid = sum(observed_data$size),
  already_paid_covid = sum(observed_data_covid$size))
```

```
## already_paid_no_covid    already_paid_covid
##                62217515                61504837
```

For **Q.2** we apply the chain ladder method to the cumulative runoff triangle.

```
require(ChainLadder)
triangle <- cumulative_triangle(observed_data_covid,
                                variable = 'size')

cl <- MackChainLadder(triangle)
cl
```

```
## MackChainLadder(Triangle = triangle)
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1  6,455,614      1.000 6,455,614         0         0      NaN
## 2  6,215,419      1.000 6,215,419         0      1,243      Inf
## 3  6,800,299      1.000 6,802,161      1,862      3,366      1.8079
## 4  6,440,919      0.999 6,444,741      3,822      5,265      1.3777
## 5  6,450,263      0.999 6,456,452      6,190      7,483      1.2089
## 6  6,323,908      0.999 6,332,152      8,243      8,128      0.9860
## 7  6,422,224      0.997 6,440,645     18,422     10,090      0.5477
## 8  6,273,540      0.991 6,330,206     56,666     22,280      0.3932
## 9  6,259,801      0.972 6,441,698     181,896     35,842      0.1970
## 10 3,862,850      0.704 5,490,850     1,628,000    124,951      0.0768
##
##              Totals
## Latest:    61,504,836.84
## Dev:              0.97
## Ultimate: 63,409,936.76
## IBNR:        1,905,099.93
## Mack.S.E     136,795.35
## CV(IBNR):      0.07
```

For **Q.3** we compute the prediction error of the chain ladder method.

```
ultimate <- sum(cum2incr(cl$FullTriangle))
already_paid <- sum(cum2incr(cl$Triangle), na.rm = TRUE)

reserve_covid_cl <- ultimate - already_paid
sigma_cl <- as.numeric(cl$Total.Mack.S.E)

error = reserve_covid_actual - reserve_covid_cl

round(c(error = error,
        pct_error = error / reserve_covid_actual * 100,
        std.dev = error / sigma_cl),2)
```

##	error	pct_error	std.dev
##	484177.34	20.26	3.54

The reduction in claim frequency in the months March, April and May has reduced the accuracy of the chain ladder method.

Actual triangle:

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	4532915	1739699	121876.31	25951.66	9304.915	6703.063	9478.770	6217.266	3468.36
##	[2,]	4414954	1673577	107415.07	19472.27	0.000	0.000	0.000	0.000	0.00
##	[3,]	4786833	1785616	149239.81	59348.18	14857.743	4404.145	0.000	0.000	0.00
##	[4,]	4653451	1614053	127900.21	36213.57	9301.109	0.000	0.000	0.000	0.00
##	[5,]	4321019	1909085	159246.74	49184.25	11727.691	0.000	8367.837	1250.830	0.00
##	[6,]	4417129	1722151	103906.89	65592.12	15129.218	0.000	0.000	0.000	0.00
##	[7,]	4538329	1718509	146138.00	19248.11	8206.603	1255.568	0.000	0.000	0.00
##	[8,]	4578285	1618201	77053.93	20586.47	4744.763	0.000	0.000	0.000	0.00
##	[9,]	4503558	1756243	153002.08	50348.72	10709.289	4574.275	0.000	0.000	0.00
##	[10,]	3862850	1801338	214885.47	77394.80	26305.781	6306.792	0.000	0.000	0.00

Predicted triangle:

##		dev								
##	origin	1	2	3	4	5	6	7	8	
##	1	4532915	1739699	121876.31	25951.66	9304.915	6703.063	9478.770	6217.266	3468.36
##	2	4414954	1673577	107415.07	19472.27	0.000	0.000	0.000	0.000	0.00
##	3	4786833	1785616	149239.81	59348.18	14857.743	4404.145	0.000	0.000	1861.5
##	4	4653451	1614053	127900.21	36213.57	9301.109	0.000	0.000	2057.632	1764.6
##	5	4321019	1909085	159246.74	49184.25	11727.691	0.000	2361.270	2061.371	1767.2
##	6	4417129	1722151	103906.89	65592.12	15129.218	2172.474	2315.810	2021.685	1733.2
##	7	4538329	1718509	146138.00	19248.11	10037.169	2209.697	2355.488	2056.324	1762.9
##	8	4578285	1618201	77053.93	38559.90	9865.059	2171.807	2315.098	2021.064	1732.7
##	9	4503558	1756243	124232.58	39239.04	10038.809	2210.058	2355.873	2056.660	1763.2
##	10	3862850	1472953	105894.83	33447.04	8556.998	1883.835	2008.127	1753.080	1502.5

When the chain ladder method fails

Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**. We construct various triangles to assess this assumption:

Number of open claims begin at the start of the year:

```
triangle_open <- incremental_triangle(  
  observed_data_covid %>%  
    mutate(open = calendar_year <= settlement_year),  
  variable = 'open')  
triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	2933	802	37	7	2	1	1	1	1	0
##	[2,]	2886	756	34	6	1	0	0	0	0	NA
##	[3,]	3000	790	41	9	5	2	0	0	NA	NA
##	[4,]	2937	791	28	8	2	1	0	NA	NA	NA
##	[5,]	2908	795	44	9	2	1	NA	NA	NA	NA
##	[6,]	2862	770	34	7	3	NA	NA	NA	NA	NA
##	[7,]	2866	788	39	7	NA	NA	NA	NA	NA	NA
##	[8,]	2862	744	25	NA	NA	NA	NA	NA	NA	NA
##	[9,]	2867	772	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	2573	NA	NA	NA	NA	NA	NA	NA	NA	NA

Settlement probability:

```
triangle_settlement <- incremental_triangle(  
  observed_data_covid %>%  
    mutate(settlement = calendar_year == settlement_year),  
  variable = 'settlement')  
  
triangle_settlement / triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	0.7565632	0.9538653	0.8108108	0.7142857	0.5	0	0	0	1	NaN
##	[2,]	0.7723493	0.9550265	0.8235294	0.8333333	1.0	NaN	NaN	NaN	NaN	NA
##	[3,]	0.7640000	0.9481013	0.7804878	0.4444444	0.6	1	NaN	NaN	NA	NA
##	[4,]	0.7647259	0.9646018	0.7142857	0.7500000	0.5	1	NaN	NA	NA	NA
##	[5,]	0.7568776	0.9446541	0.7954545	0.7777778	0.5	0	NA	NA	NA	NA
##	[6,]	0.7641509	0.9558442	0.7941176	0.5714286	1.0	NA	NA	NA	NA	NA
##	[7,]	0.7616888	0.9505076	0.8205128	0.5714286	NA	NA	NA	NA	NA	NA
##	[8,]	0.7672956	0.9663978	0.7600000	NA	NA	NA	NA	NA	NA	NA
##	[9,]	0.7697942	0.9572539	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	0.7298873	NA	NA	NA	NA	NA	NA	NA	NA	NA

The percentage of settled claims is lower in the last accident year. Indicating that more claims are still open.

Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

We construct various triangles to assess this assumption:

Percentage of open claims for which a payment is made.

```
triangle_payment <- incremental_triangle(  
  observed_data_covid,  
  variable = 'payment')
```

```
triangle_payment / triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	0.7719059	0.6857855	0.5405405	0.7142857	1.0	1.0	1	1	1	NaN
##	[2,]	0.7862093	0.7050265	0.7941176	0.8333333	0.0	NaN	NaN	NaN	NaN	NA
##	[3,]	0.7690000	0.7050633	0.6341463	0.8888889	0.6	0.5	NaN	NaN	NA	NA
##	[4,]	0.7793667	0.6788875	0.8214286	1.0000000	0.5	0.0	NaN	NA	NA	NA
##	[5,]	0.7699450	0.7308176	0.7500000	0.8888889	0.5	0.0	NA	NA	NA	NA
##	[6,]	0.7725367	0.6909091	0.7058824	0.8571429	1.0	NA	NA	NA	NA	NA
##	[7,]	0.7665736	0.6814721	0.7692308	0.7142857	NA	NA	NA	NA	NA	NA
##	[8,]	0.7861635	0.6854839	0.7200000	NA	NA	NA	NA	NA	NA	NA
##	[9,]	0.7921172	0.7046632	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	0.7765255	NA	NA	NA	NA	NA	NA	NA	NA	NA

Average amount paid given a payment is made.

```
triangle_size <- incremental_triangle(  
  observed_data_covid,  
  variable = 'size')
```

```
triangle_size / triangle_payment
```

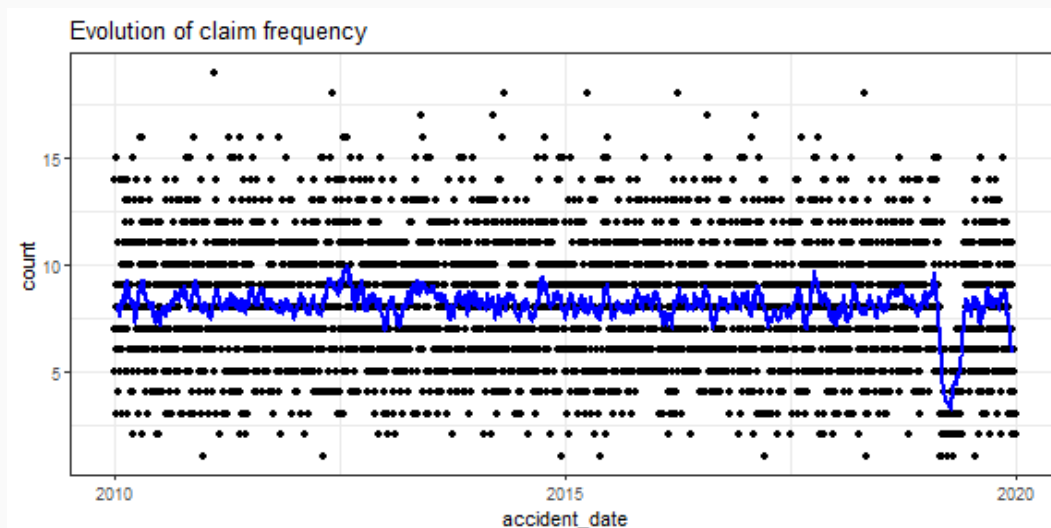
##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	2002.171	3163.090	6093.815	5190.332	4652.457	6703.063	9478.77	6217.266	3468.3	NA
##	[2,]	1945.771	3139.920	3978.336	3894.454	NaN	NaN	NaN	NaN	NA	NA
##	[3,]	2074.917	3205.774	5739.993	7418.522	4952.581	4404.145	NaN	NaN	NA	NA
##	[4,]	2032.962	3005.686	5560.879	4526.696	9301.109	NaN	NaN	NA	NA	NA
##	[5,]	1929.888	3285.860	4825.659	6148.031	11727.691	NaN	NA	NA	NA	NA
##	[6,]	1997.797	3237.126	4329.454	10932.020	5043.073	NA	NA	NA	NA	NA
##	[7,]	2065.694	3200.202	4871.267	3849.622	NA	NA	NA	NA	NA	NA
##	[8,]	2034.793	3172.943	4280.774	NA	NA	NA	NA	NA	NA	NA
##	[9,]	1983.073	3228.389	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	1933.359	NA	NA	NA	NA	NA	NA	NA	NA	NA

Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

Inspecting the (granular) daily data:

```
claims_covid <- observed_data_covid %>%  
  group_by(accident_number) %>%  
  slice(1) %>%  
  ungroup()  
  
occ_intensity <- claims_covid %>%  
  group_by(accident_date) %>%  
  summarise(count = n())  
  
require(zoo)  
occ_intensity$moving_average <-  
  rollmean(occ_intensity$count, 30, na.pad = TRUE)  
  
ggplot(occ_intensity) +  
  theme_bw() +  
  geom_point(aes(x = accident_date, y = count)) +  
  geom_line(aes(x = accident_date, y = moving_average),  
            size = 1, color = 'blue') +  
  ggtitle('Evolution of claim frequency')
```



The moving average indicates a period with decreased claim frequency in 2019

Detecting when the chain ladder method fails

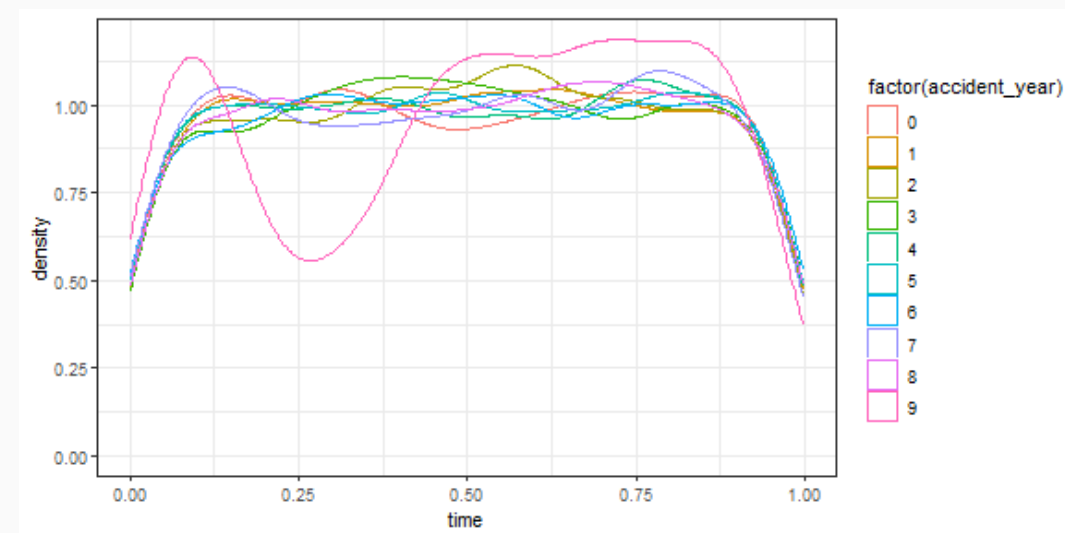
The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

Inspecting the (granular) daily data:

```
require(lubridate)
claims_covid <- claims_covid %>%
  mutate(start_year = floor_date(accident_date, unit = 'year'),
         time = as.numeric(accident_date - start_year) / 366)

ggplot(claims_covid) +
  theme_bw() +
  geom_density(aes(x = time,
                  group = factor(accident_year),
                  color = factor(accident_year)))
```

Density of when claims occur within each accident year.
The chain ladder method works best when the densities look similar for all accident years.



The occurrence intensity in accident year 9 deviates from the other accident years.

Fixing the chain ladder method

Improve the performance of the chain ladder method by making the **data more homogeneous**:

- **Group** claims by:
 - claim characteristics, e.g. separate triangle for fire and water related claims in home insurance.
 - occurrence time within a calendar year
- **Omit**:
 - deviating years/cells, especially if they occurred long ago
- **Adjust** for:
 - inflation

The appropriate method(s) depend on the data inspection from the previous slides.

Fixing the chain ladder method (continued)

Applied to the **covid data set**:

- claim frequency has changed in the months March, April, May
- only 2019, accident year 9, is affected

As 2019 is a recent calendar year, we cannot omit accident year 9 from the data.

We improve homogeneity in the data by splitting claims in two groups:

- group A: claims that have occurred in March, April or May
- group B: claims that have not occurred in March, April or May

```
observed_data_covid <- observed_data_covid %>%  
  mutate(occurrence_month = as.numeric(format(accident_date, '%m')))  
  
observed_data_covidA <- observed_data_covid %>%  
  filter(occurrence_month %in% c(3, 4, 5))  
  
observed_data_covidB <- observed_data_covid %>%  
  filter(!(occurrence_month %in% c(3, 4, 5)))
```



Your turn

In this exercise you compute the chain ladder reserve for claims in group A and B.

1. Choose data set A or B and analyse the stability of the chosen data set using some of the presented detection tools.
2. Compute the chain ladder reserve for group A and B separately.
Combine these estimates to obtain an estimate for the total reserve.
3. Compare the new reserve estimate with the estimate obtained earlier using the chain ladder method without splitting the data.
4. Compute the standard error of the reserve estimate using the formula

$$\sigma^2 = \sigma_A^2 + \sigma_B^2.$$

For **Q1**, analysing the stability of data set A.

Number of open claims begin dev.year:

```
triangle_open <- incremental_triangle(  
  observed_data_covidA %>%  
    mutate(open = calendar_year <= settlement_year),  
  variable = 'open')  
triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	775	53	7	1	1	0	0	0	0	0
##	[2,]	748	38	0	0	0	0	0	0	0	NA
##	[3,]	773	47	5	1	1	0	0	0	NA	NA
##	[4,]	781	52	3	0	0	0	0	NA	NA	NA
##	[5,]	757	63	7	2	1	0	NA	NA	NA	NA
##	[6,]	730	38	5	1	1	NA	NA	NA	NA	NA
##	[7,]	750	59	2	1	NA	NA	NA	NA	NA	NA
##	[8,]	704	51	8	NA	NA	NA	NA	NA	NA	NA
##	[9,]	770	40	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	367	NA	NA	NA	NA	NA	NA	NA	NA	NA

Settlement probability:

```
triangle_settlement <- incremental_triangle(  
  observed_data_covidA %>%  
    mutate(settlement = calendar_year == settlement_year),  
  variable = 'settlement')  
  
triangle_settlement / triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	0.9316129	0.8679245	0.8571429	0.0	1	NaN	NaN	NaN	NaN	NaN
##	[2,]	0.9491979	1.0000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NA
##	[3,]	0.9391979	0.8936170	0.8000000	0.0	1	NaN	NaN	NaN	NA	NA
##	[4,]	0.9334187	0.9423077	1.0000000	NaN	NaN	NaN	NaN	NA	NA	NA
##	[5,]	0.9167768	0.8888889	0.7142857	0.5	1	NaN	NA	NA	NA	NA
##	[6,]	0.9479452	0.8684211	0.8000000	0.0	1	NA	NA	NA	NA	NA
##	[7,]	0.9213333	0.9661017	0.5000000	0.0	NA	NA	NA	NA	NA	NA
##	[8,]	0.9275568	0.8431373	0.6250000	NA	NA	NA	NA	NA	NA	NA
##	[9,]	0.9480519	0.8750000	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	0.9455041	NA	NA	NA	NA	NA	NA	NA	NA	NA

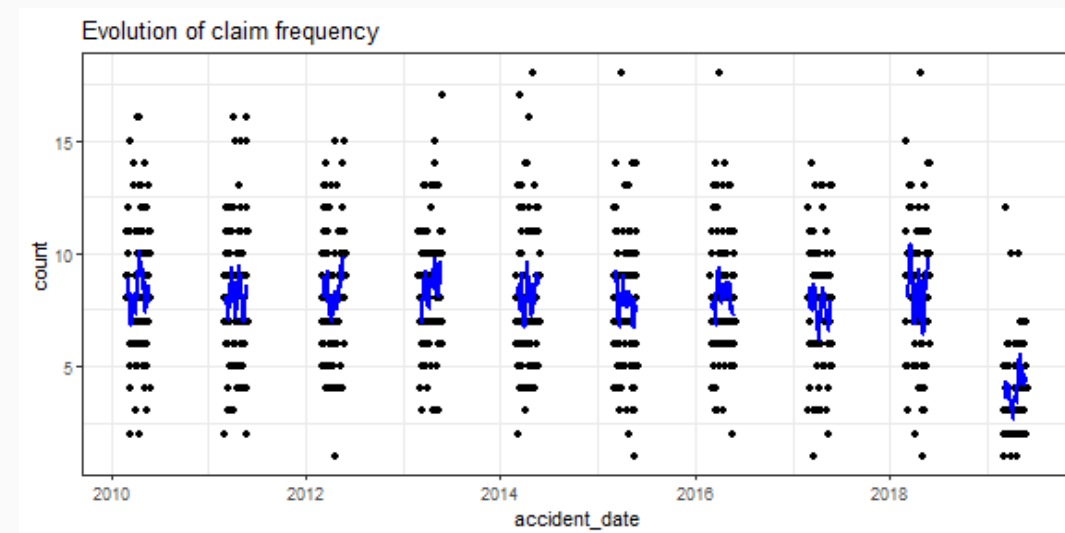
For **Q.1**, analysing the stability of data set A.

```
claims_covidA <- observed_data_covidA %>%
  group_by(accident_number) %>%
  slice(1) %>%
  ungroup()

occ_intensity <- claims_covidA %>%
  group_by(accident_date, accident_year) %>%
  summarise(count = n())

require(zoo)
occ_intensity <- occ_intensity %>%
  group_by(accident_year) %>%
  mutate(moving_average = rollmean(count, 14, na.pad = TRUE))

ggplot(occ_intensity) +
  theme_bw() +
  geom_point(aes(x = accident_date, y = count)) +
  geom_line(aes(x = accident_date, y = moving_average),
            size = 1, color = 'blue') +
  ggtitle('Evolution of claim frequency')
```

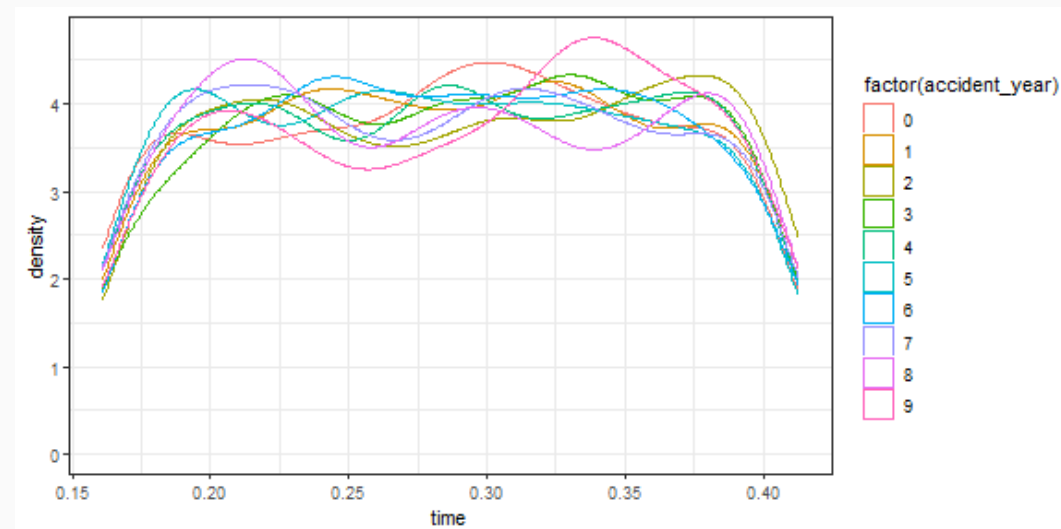


We compute moving averages within each accident year. Otherwise `rollmean` would naively assume that May 31, 2010 is followed by March 1, 2011.

For **Q.1**, analysing the stability of data set A.

```
require(lubridate)
claims_covidA <- claims_covidA %>%
  mutate(start_year = floor_date(accident_date, unit = 'year'),
         time = as.numeric(accident_date - start_year) / 366)

ggplot(claims_covidA) +
  theme_bw() +
  geom_density(aes(x = time,
                  group = factor(accident_year),
                  color = factor(accident_year)))
```



All densities have a similar, indicating that within each accident year the distributon of when claims occur is the same.

For **Q.1**, analysing the stability of data set B.

Number of open claims begin at the start of the year:

```
triangle_open <- incremental_triangle(  
  observed_data_covidB %>%  
    mutate(open = calendar_year <= settlement_year),  
  variable = 'open')  
triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	2158	749	30	6	1	1	1	1	1	0
##	[2,]	2138	718	34	6	1	0	0	0	0	NA
##	[3,]	2227	743	36	8	4	2	0	0	NA	NA
##	[4,]	2156	739	25	8	2	1	0	NA	NA	NA
##	[5,]	2151	732	37	7	1	1	NA	NA	NA	NA
##	[6,]	2132	732	29	6	2	NA	NA	NA	NA	NA
##	[7,]	2116	729	37	6	NA	NA	NA	NA	NA	NA
##	[8,]	2158	693	17	NA	NA	NA	NA	NA	NA	NA
##	[9,]	2097	732	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	2206	NA	NA	NA	NA	NA	NA	NA	NA	NA

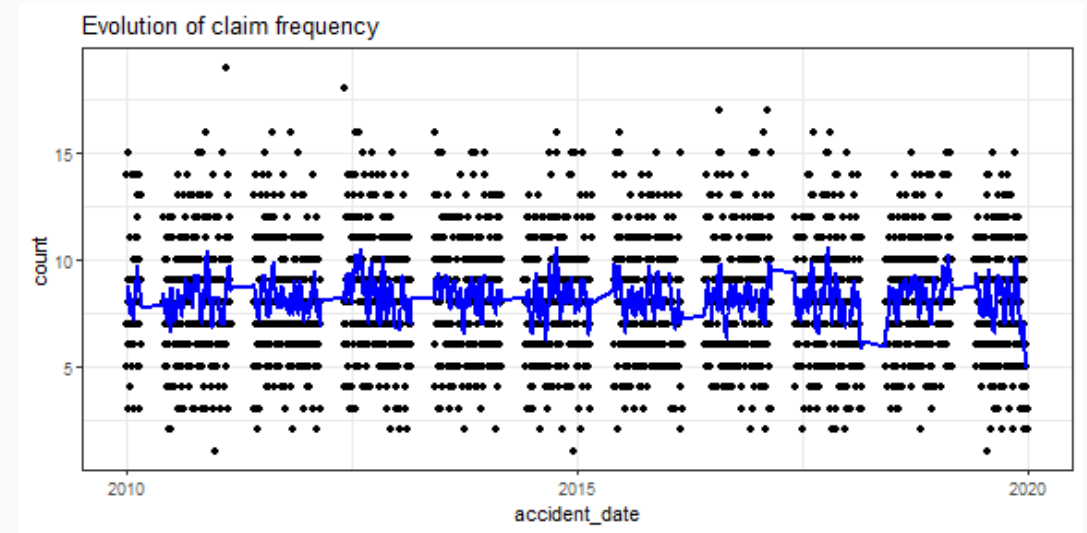
Settlement probability:

```
triangle_settlement <- incremental_triangle(  
  observed_data_covidB %>%  
    mutate(settlement = calendar_year == settlement_year),  
  variable = 'settlement')  
  
triangle_settlement / triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	0.6936979	0.9599466	0.8000000	0.8333333	0.0	0	0	0	1	NaN
##	[2,]	0.7104771	0.9526462	0.8235294	0.8333333	1.0	NaN	NaN	NaN	NaN	NA
##	[3,]	0.7031881	0.9515478	0.7777778	0.5000000	0.5	1	NaN	NaN	NA	NA
##	[4,]	0.7036178	0.9661705	0.6800000	0.7500000	0.5	1	NaN	NA	NA	NA
##	[5,]	0.7006044	0.9494536	0.8108108	0.8571429	0.0	0	NA	NA	NA	NA
##	[6,]	0.7012195	0.9603825	0.7931034	0.6666667	1.0	NA	NA	NA	NA	NA
##	[7,]	0.7051040	0.9492455	0.8378378	0.6666667	NA	NA	NA	NA	NA	NA
##	[8,]	0.7150139	0.9754690	0.8235294	NA	NA	NA	NA	NA	NA	NA
##	[9,]	0.7043395	0.9617486	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	0.6940163	NA	NA	NA	NA	NA	NA	NA	NA	NA

For **Q.1**, analysing the stability of data set B.

```
claims_covidB <- observed_data_covidB %>%  
  group_by(accident_number) %>%  
  slice(1) %>%  
  ungroup()  
  
occ_intensity <- claims_covidB %>%  
  group_by(accident_date, accident_year) %>%  
  summarise(count = n())  
  
require(zoo)  
occ_intensity <- occ_intensity %>%  
  group_by(accident_year) %>%  
  mutate(moving_average = rollmean(count, 14, na.pad = TRUE))  
  
ggplot(occ_intensity) +  
  theme_bw() +  
  geom_point(aes(x = accident_date, y = count)) +  
  geom_line(aes(x = accident_date, y = moving_average),  
            size = 1, color = 'blue') +  
  ggtitle('Evolution of claim frequency')
```



Claim frequency is stable over all accident years. No sudden jumps/decreases.

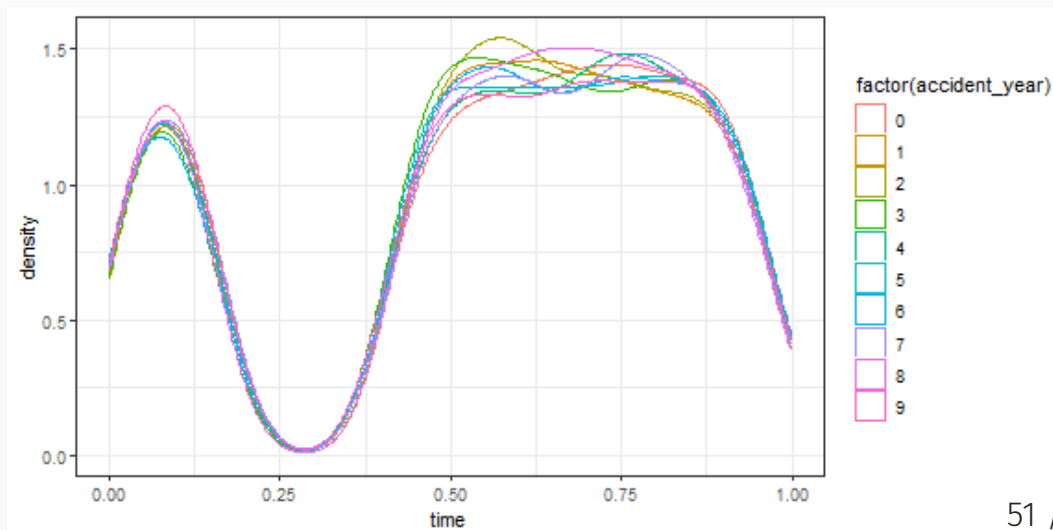
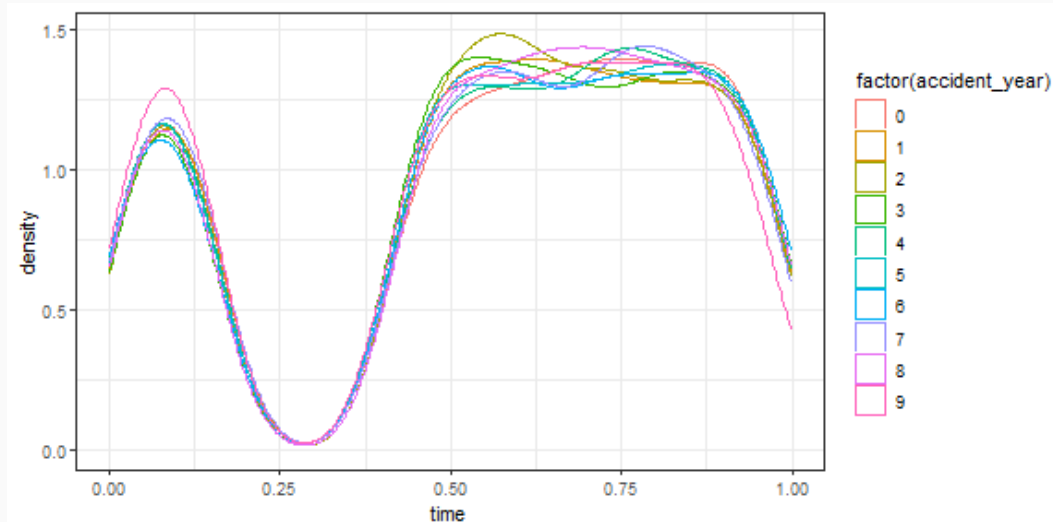
For **Q.1**, analysing the stability of data set B.

```
require(lubridate)
claims_covidB <- claims_covidB %>%
  mutate(start_year = floor_date(accident_date, unit = 'year'),
         time = as.numeric(accident_date - start_year) / 366)

ggplot(claims_covidB) +
  theme_bw() +
  geom_density(aes(x = time,
                  group = factor(accident_year),
                  color = factor(accident_year)))
```

Since we only observe reported claims, we observe less claims from recent accident years around the end of the year. Correct for this by filtering claims that are reported in the year of occurrence.

```
ggplot(claims_covidB %>%
  filter(reporting_year == accident_year)) +
  theme_bw() +
  geom_density(aes(x = time,
                  group = factor(accident_year),
                  color = factor(accident_year)))
```



For **Q.2** computing the chain ladder reserve

```
require(ChainLadder)
triangleA <- cumulative_triangle(observed_data_covidA,
                                variable = 'size')
clA <- MackChainLadder(triangleA)
clA
```

```
## Warning in Mack.S.E(CL[["Models"]], FullTriangle, est.sigma = est.sigma, : Informatio
## '5-6', '6-7', '7-8', '8-9'
```

```
## MackChainLadder(Triangle = triangleA)
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1  1,619,982      1.000 1,619,982  0.00e+00 0.00e+00      NaN
## 2  1,528,185      1.000 1,528,185  0.00e+00 1.44e-13      Inf
## 3  1,608,274      1.000 1,608,274 -4.66e-10 3.54e-10    -0.761
## 4  1,671,348      1.000 1,671,348 -4.66e-10 3.64e-10    -0.781
## 5  1,684,656      1.000 1,684,656 -4.66e-10 3.66e-10    -0.785
## 6  1,470,984      1.000 1,470,984 -4.66e-10 4.23e-10    -0.908
## 7  1,640,175      0.999 1,641,391  1.22e+03 2.38e+03    1.962
## 8  1,666,274      0.996 1,672,544  6.27e+03 4.74e+03    0.756
## 9  1,615,375      0.987 1,636,318  2.09e+04 1.18e+04    0.561
## 10 724,846       0.893  811,877  8.70e+04 3.14e+04    0.361
##
##              Totals
## Latest:   15,230,099.47
## Dev:             0.99
## Ultimate: 15,345,559.70
## IBNR:        115,460.23
## Mack.S.E      34,373.24
## CV(IBNR):         0.30
```

```
require(ChainLadder)
triangleB <- cumulative_triangle(observed_data_covidB,
                                variable = 'size')
clB <- MackChainLadder(triangleB)
clB
```

```
## MackChainLadder(Triangle = triangleB)
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1  4,835,633      1.000 4,835,633      0      0      NaN
## 2  4,687,234      1.000 4,687,234      0  1,212      Inf
## 3  5,192,025      1.000 5,193,916  1,892  3,398  1.7965
## 4  4,769,571      0.999 4,773,326  3,755  5,225  1.3914
## 5  4,765,606      0.999 4,771,681  6,074  7,422  1.2219
## 6  4,852,924      0.998 4,861,338  8,414  8,249  0.9803
## 7  4,782,048      0.996 4,799,125 17,077  9,690  0.5675
## 8  4,607,266      0.989 4,656,956 49,690 19,630  0.3950
## 9  4,644,426      0.967 4,804,801 160,374 39,868  0.2486
## 10 3,138,005      0.640 4,902,249 1,764,244 132,487  0.0751
##
##              Totals
## Latest:   46,274,737.36
## Dev:             0.96
## Ultimate: 48,286,257.75
## IBNR:        2,011,520.39
## Mack.S.E      144,671.48
## CV(IBNR):         0.07
```

For **Q.3**, we combine the estimated reserves from A and B

```
ultimateA <- sum(cum2incr(clA$FullTriangle))
already_paidA <- sum(cum2incr(clA$Triangle), na.rm = TRUE)
reserveA <- ultimateA - already_paidA

ultimateB <- sum(cum2incr(clB$FullTriangle))
already_paidB <- sum(cum2incr(clB$Triangle), na.rm = TRUE)
reserveB <- ultimateB - already_paidB

reserve_covid_cl_split <- reserveA + reserveB

c(reserve_actual = reserve_covid_actual,
  reserve_cl = reserve_covid_cl,
  reserve_cl_split = reserve_covid_cl_split)
```

```
##   reserve_actual   reserve_cl reserve_cl_split
##           2389277         1905100         2126981
```

Splitting the data results in a more accurate reserve estimate.

For **Q.4**, we evaluate the relative performance of the model

```
sigma = as.numeric(
  sqrt(clA$Total.Mack.S.E^2 + clB$Total.Mack.S.E^2))

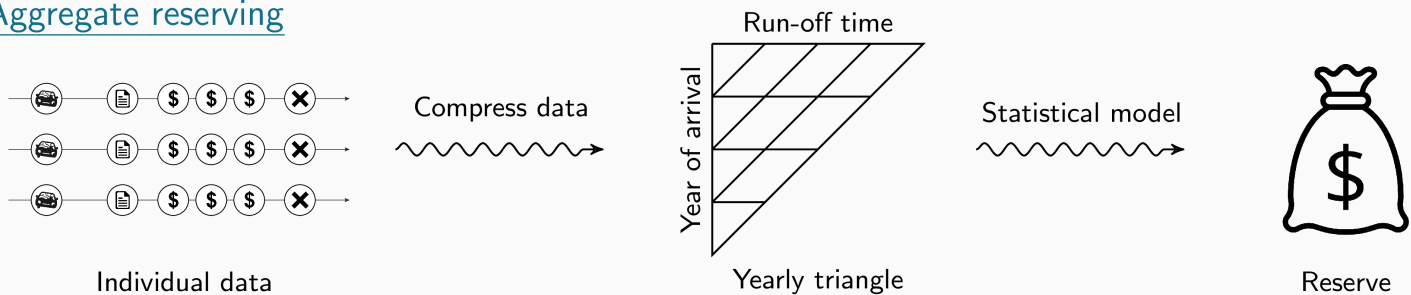
data.frame(
  method = c('chain ladder', 'chain ladder split'),
  error = c(reserve_covid_cl - reserve_covid_actual,
            reserve_covid_cl_split - reserve_covid_actual),
  sigma = c(sigma_cl, sigma)) %>%
  mutate(pct_error = error / reserve_covid_actual * 100,
         std.dev = error / sigma)
```

	chain ladder	chain ladder split
error	-484177.34	-262296.65
sigma	136795.35	148698.88
pct_error	-20.26	-10.98
std.dev	-3.54	-1.76

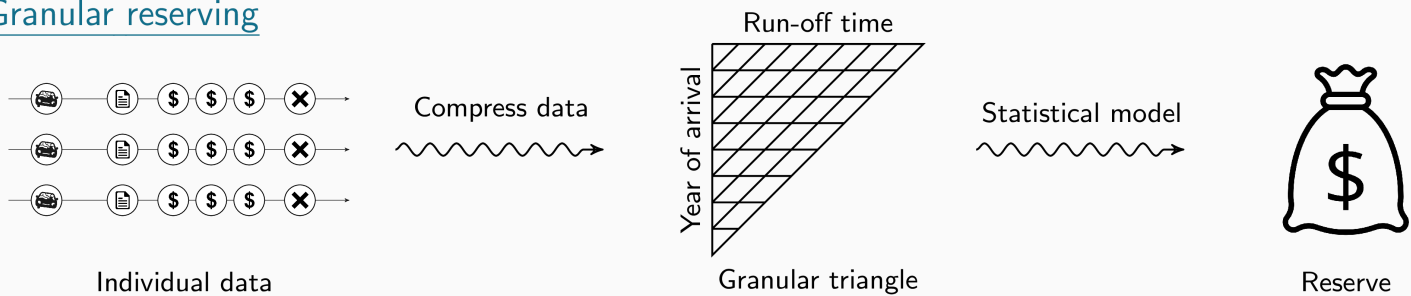
Research outlook

Research outlook

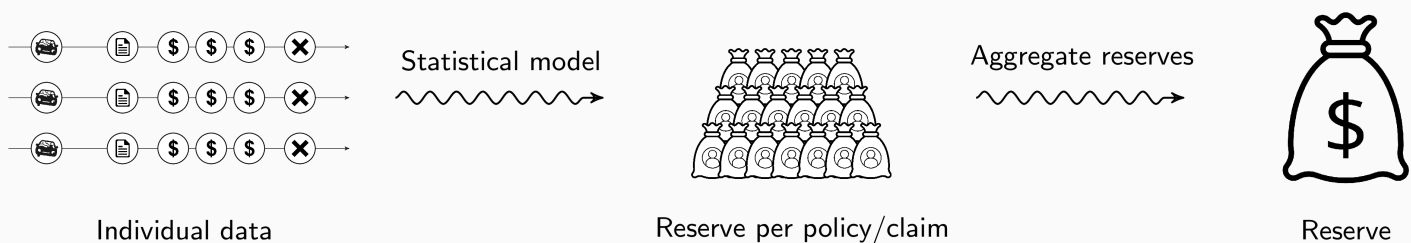
Aggregate reserving



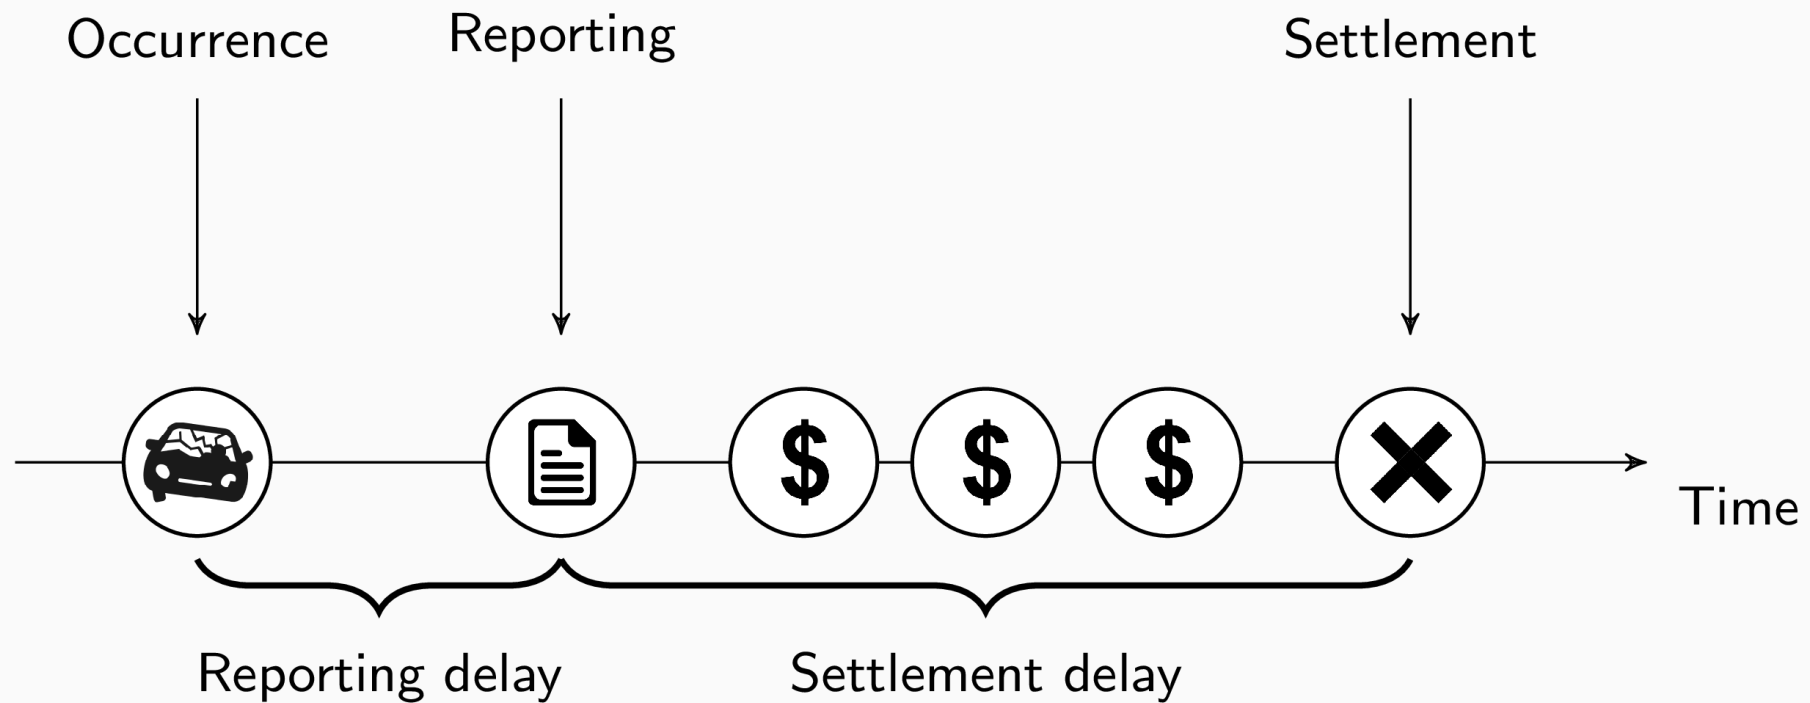
Granular reserving



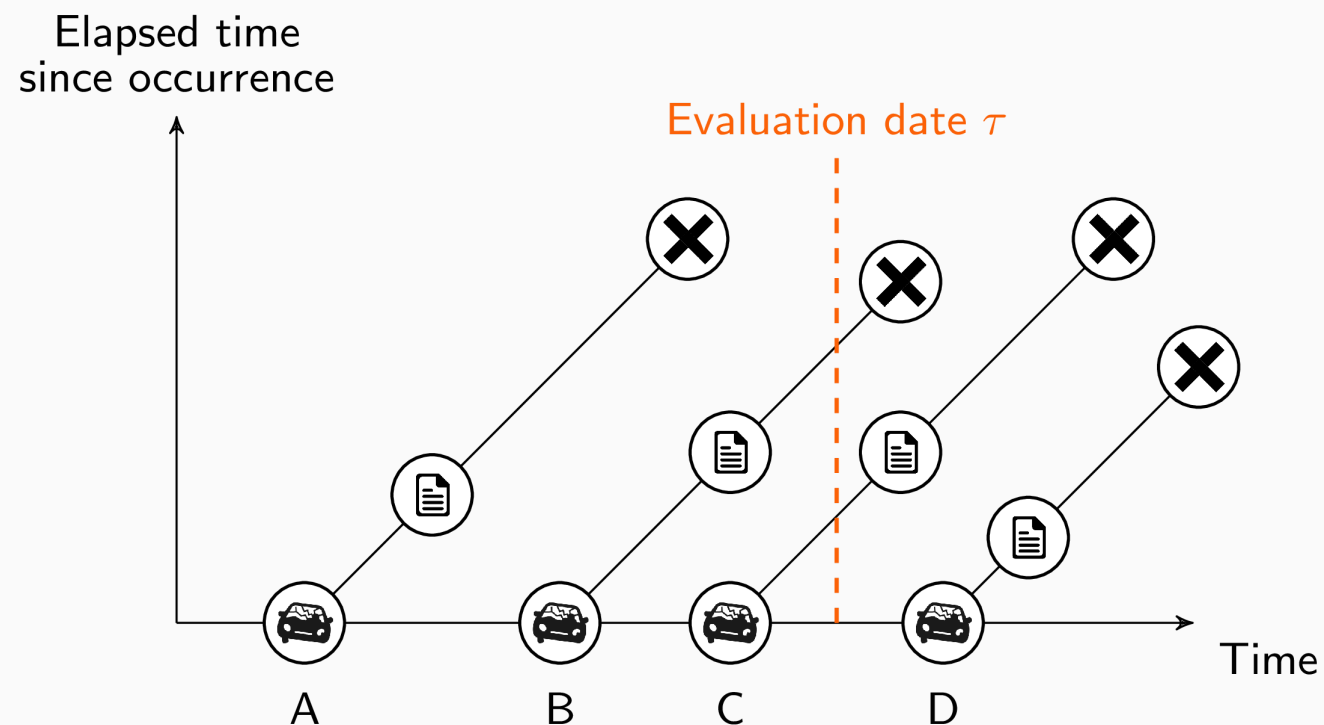
Individual reserving



Development of a individual claims



Development of a individual claims (continued)



Observed claims are **censored** due to **delays** (reporting, settlement) in the claim development process

IBNR reserve: future costs for claims that occurred, but are not yet reported (claim B)

RBNS reserve: future costs for claims that are reported, but are not yet settled (claim C)

Pricing: all costs for claims that will occur in **future** insured **exposure periods** (claim D)

The IBNR reserve

Following a frequency-severity decomposition, the IBNR reserve is the product of the expected **number of unreported claims** times the **expected severity per claim**.

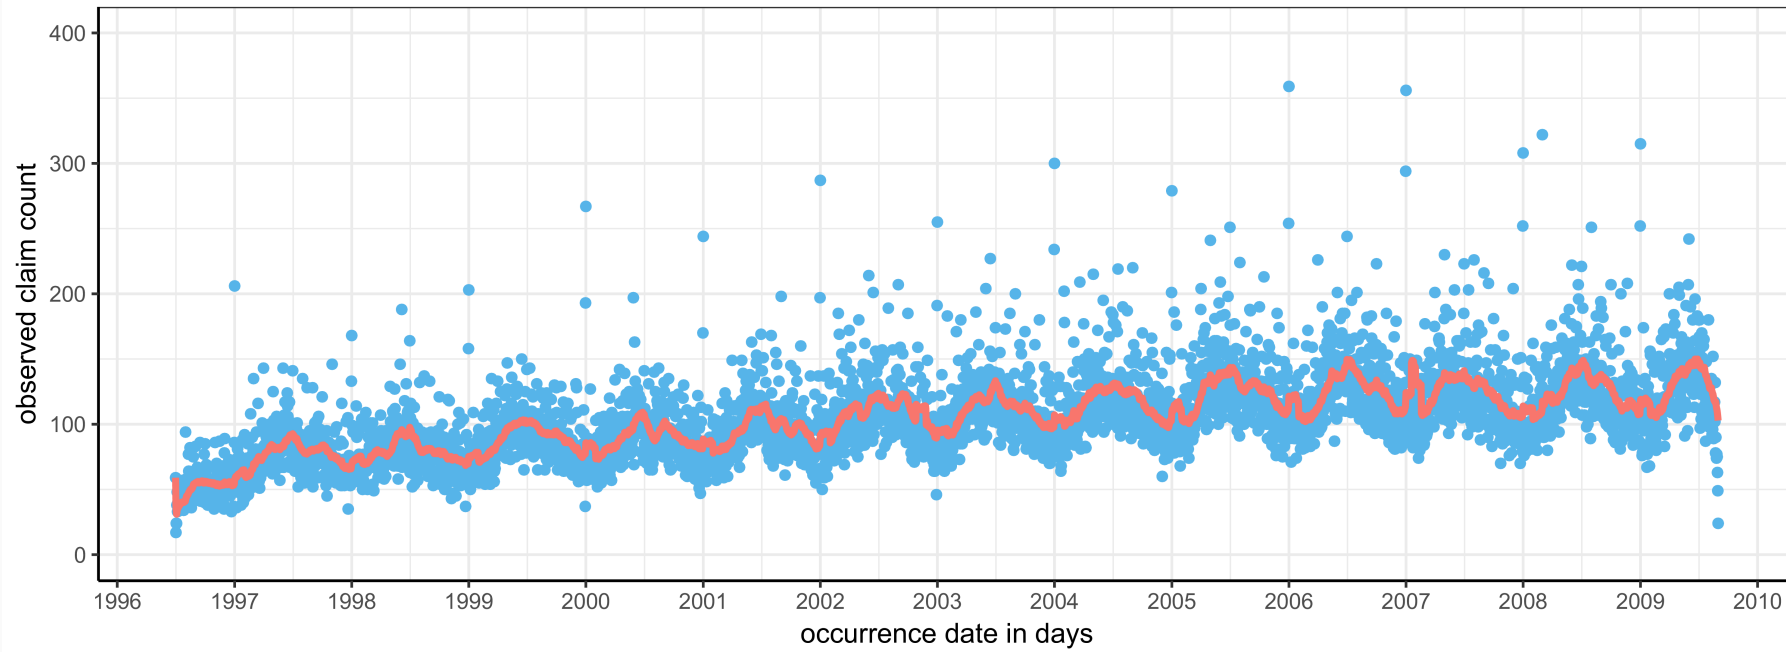
- Insurance pricing covers the estimation of claim severity for new claims
- Reserving methods focus on predicting the **number of unreported claims**.

Our work on the topic:

- "Modeling the number of hidden events subject to observation delay". Jonas Crevecoeur, Katrien Antonio and Roel Verbelen. (2019). European Journal of Operational Research.
- "Modeling the occurrence of events subject to a reporting delay via an EM algorithm". Roel Verbelen, Katrien Antonio, Gerda Claeskens and Jonas Crevecoeur. (2019). Submitted.

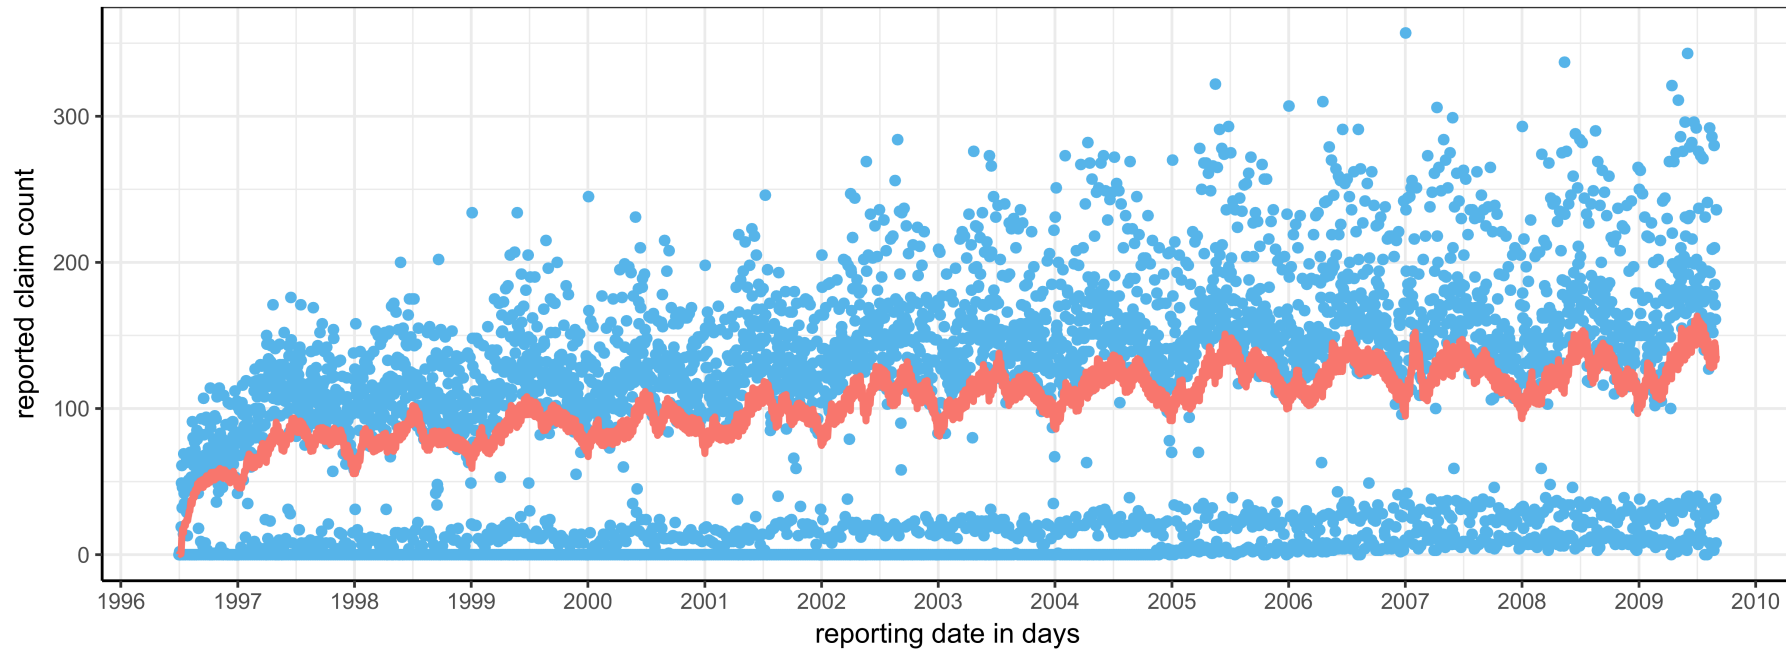
The IBNR reserve (continued)

Number of observed claim occurrences per day



The IBNR reserve (continued)

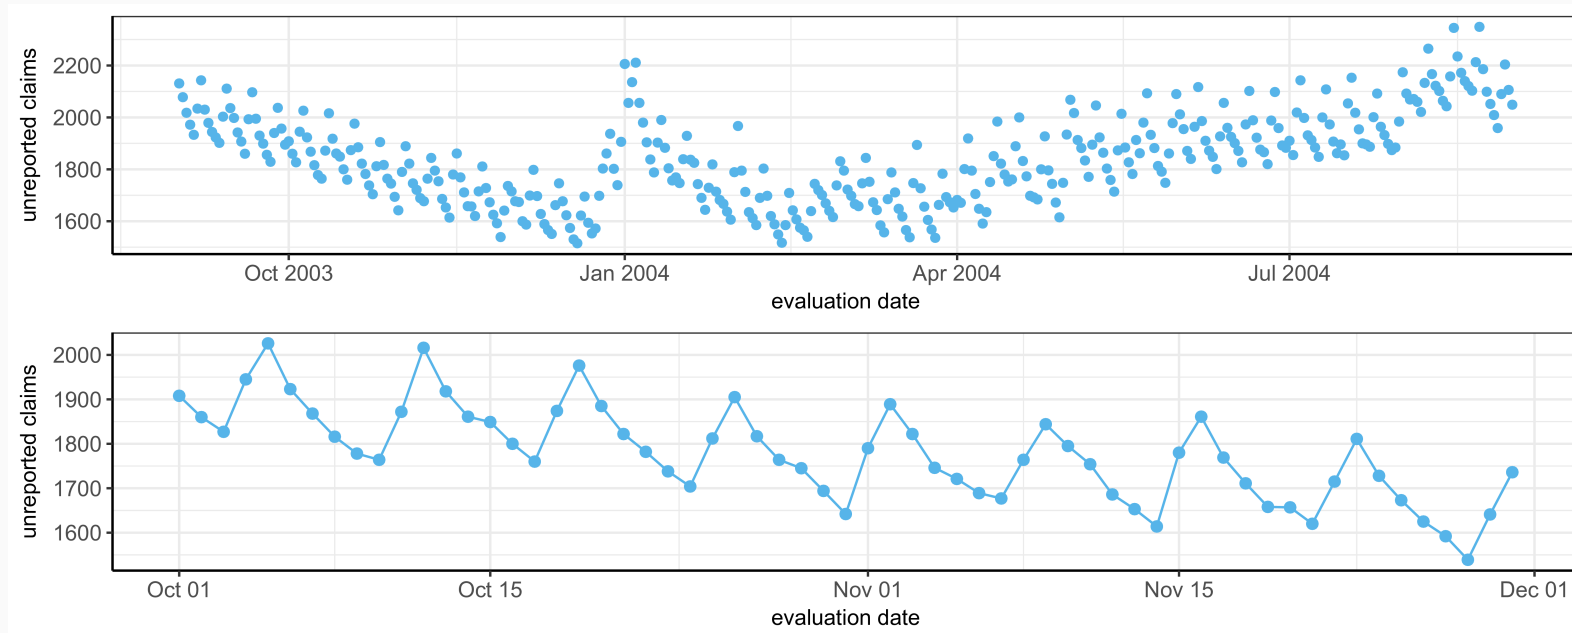
Number of reported claims per day



Almost no claims are reported on Saturdays, Sundays and holidays.

The IBNR reserve (continued)

Total number of unreported claims on each day, i.e. the number of claims that occur before the evaluation date, but are reported afterwards.



The number of unreported claims increases during the weekend (+10% on sunday) and around the end of the year (+30%).

The RBNS reserve

Predict the future costs of individual, reported claims.

Our work on the topic:

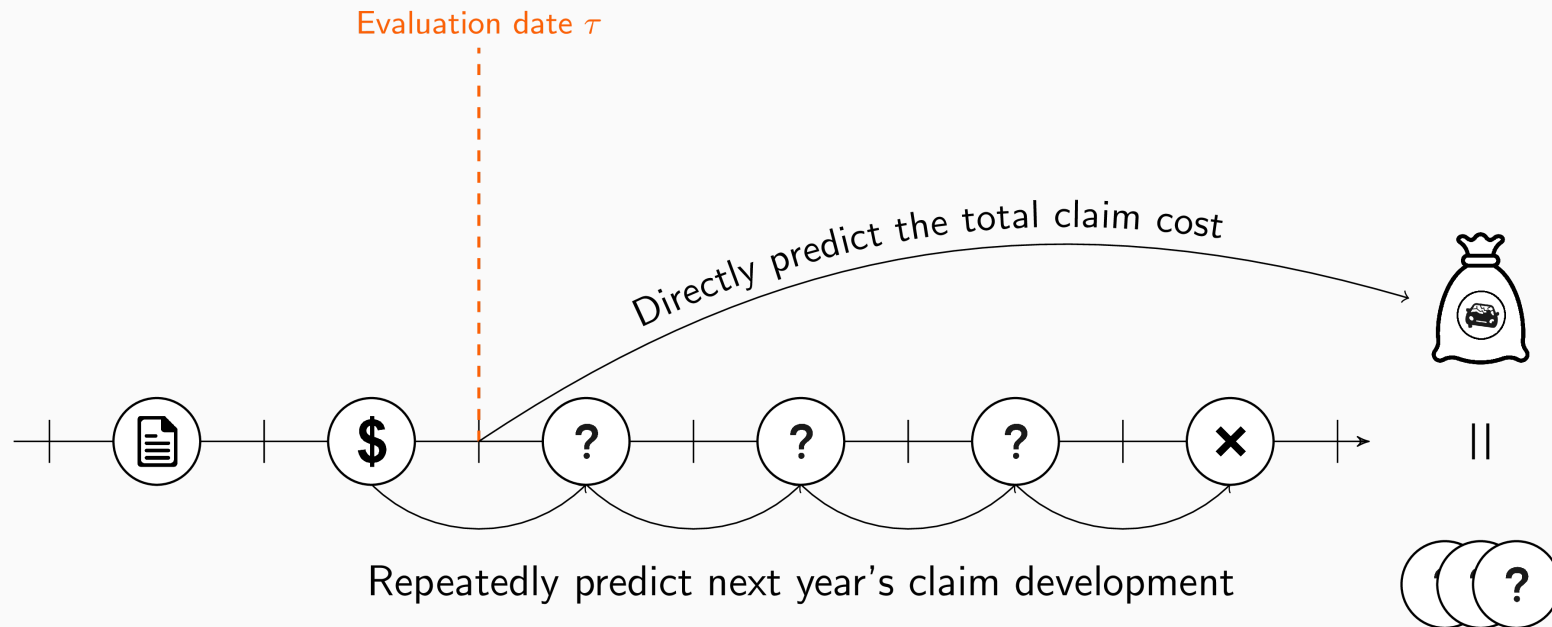
- "A hierarchical reserving model for non-life insurance claims". Jonas Crevecoeur and Katrien Antonio. (2020). Sumbitted.
- "Bridging the gap between pricing and reserving with an occurrence and development model for non-life insurance claims". Katrien Antonio and Jonas Crevecoeur. (2020). Working paper.

Methods implemented in the R package {hirem}, Hierarchical reserving models, available on [GitHub](#).

The RBNS reserve (continued)

The hierarchical model is based on **two key ideas**:

1. Model the development of individual claim in discrete time (steps of one year) using the observed development in previous years.



The RBNS reserve (continued)

The hierarchical model is based on **two key ideas**:

1. Model the development of individual claim in discrete time (steps of one year) using the observed development in previous years.
2. Focus on all events registered over the lifetime of a claim

Common events registered over the lifetime of a claim:

- settlement
- payments
- initial incurred / changes in the incurred
- involvement lawyer

These events are **dependent**:

- if a claim is **settled**, there will be no **payments** in the future
- large **payments** are more likely when the outstanding **reserve** is large

Implementation in the {hirem} package

Implementation readily available from the {hirem} package on [GitHub](#).

Events are added to the model as **layers**.

```
require(hirem)

individual_data <- individual_data %>%
  mutate(calendar_year = accident_year + development_year - 1,
         close = (settlement_year == calendar_year))

model <- hirem(individual_data %>% filter(calendar_year <= settlement_year)) %>%
  layer_glm('close', binomial(link = logit)) %>%
  layer_glm('payment', binomial(link = logit)) %>%
  layer_glm('size', Gamma(link = log),
         filter = function(data){data$payment == 1})

model <- fit(model,
  close = 'close ~ factor(development_year)',
  payment = 'payment ~ close + factor(development_year)',
  size = 'size ~ close + factor(development_year)')
```

Thanks!

Slides created with the R package `xaringan`.

Course material available via

 <https://github.com/katrienantonio/workshop-loss-reserv-fraud-2020>