

# Loss modelling and reserving analytics in R

A hands-on workshop

---

Katrien Antonio & Jonas Crevecoeur

IA|BE workshop | June 3, 10 & 17, 2021

# Today's Outline

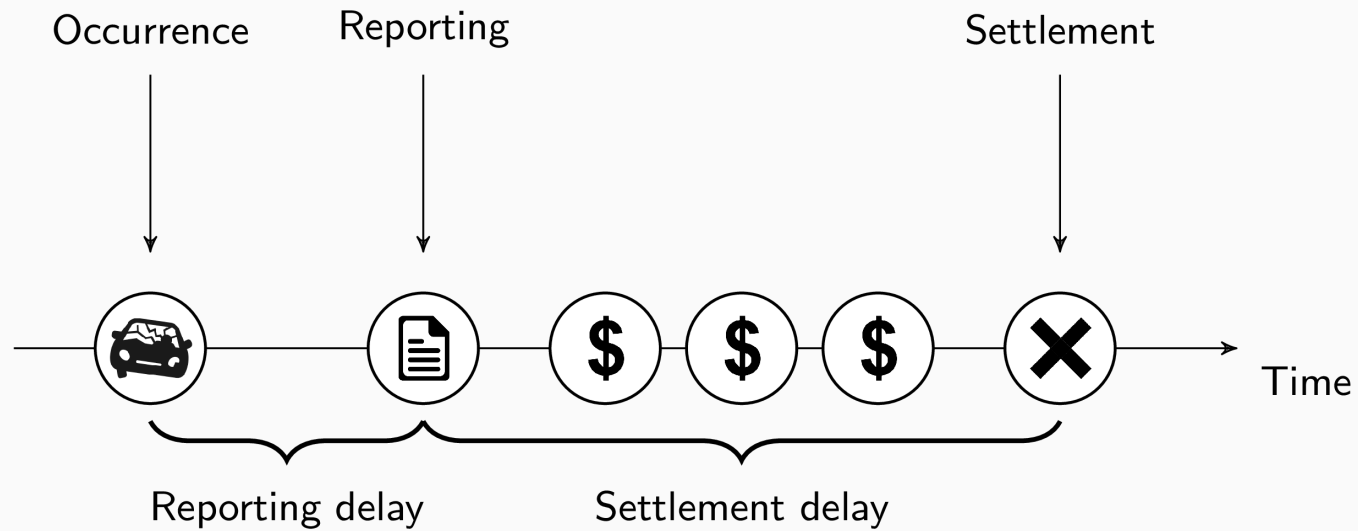
- Motivation and strategies
- Reserving data structures
  - individual and aggregated reserving data
  - runoff triangles
- Claims reserving with triangles
  - chain ladder method
  - {ChainLadder} package
  - chain ladder implementation with GLMs
- When the chain ladder method fails
  - detection tools for triangle stability
  - creating homogeneous triangles
- Research outlook
  - granular reserving methods for predicting the IBNR reserve
  - individual reserving methods for predicting the RBNS reserve

# Motivation and strategies

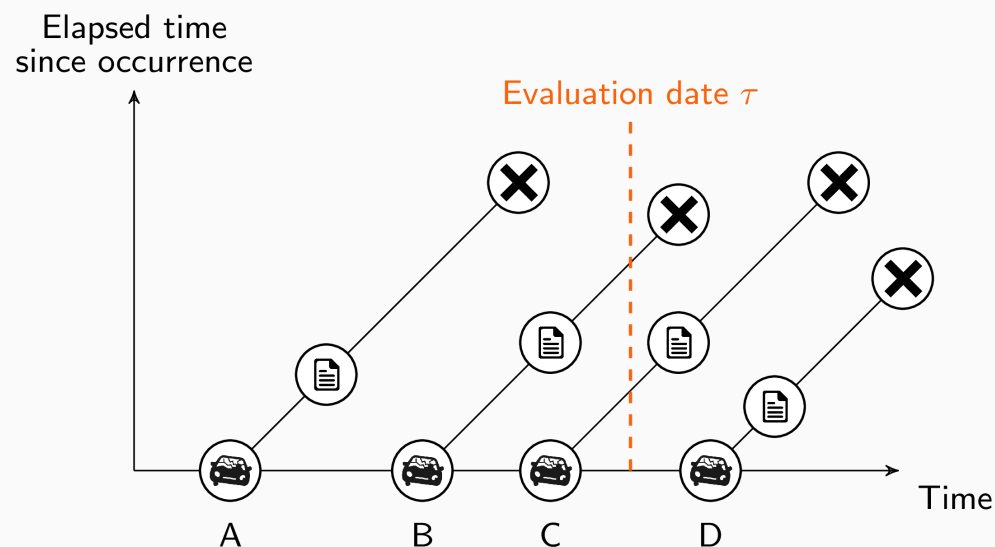
---

# Motivation

Insurers observe the **detailed evolution** of **individual claims** over time:



# Motivation (continued)



Observed claims are **censored** due to **delays** (reporting, settlement) in the claim development process.

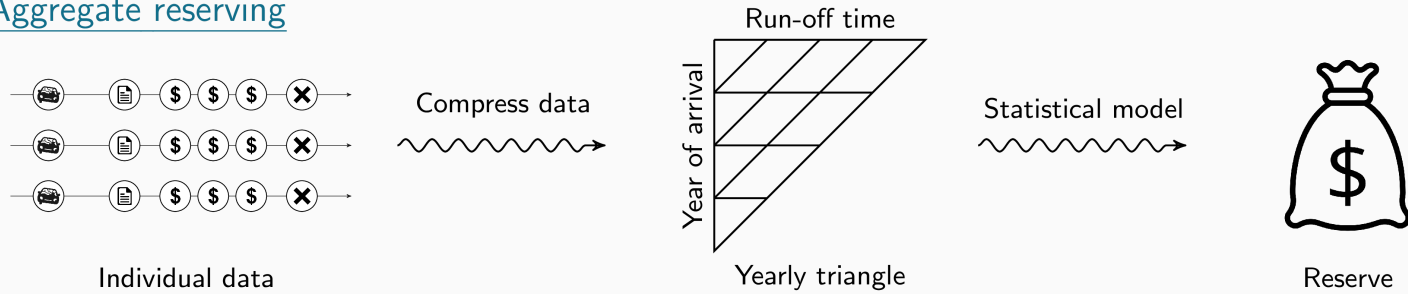
**Reserve**: future costs for claims that occurred in **past exposure periods** (claims B and C)

- RBNS: Reported, But Not yet Settled claims (claim B)
- IBNR: Incurred, But Not yet Reported claims (claim C).

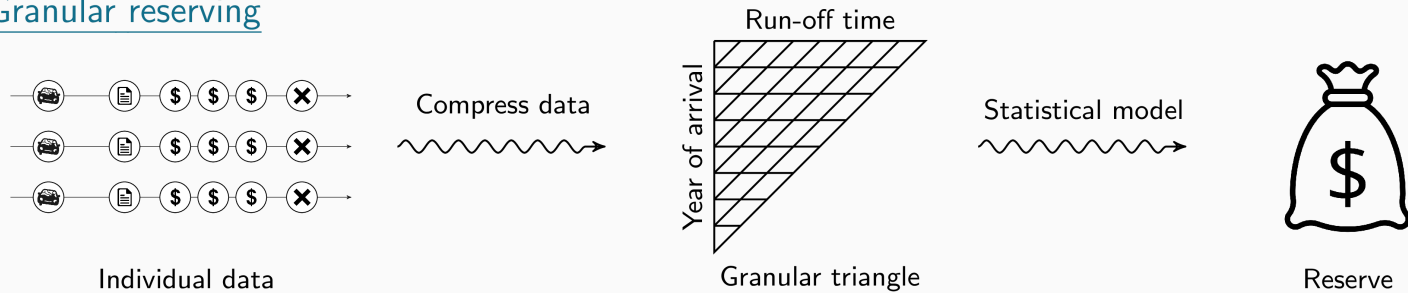
**Pricing**: all costs for claims that will occur in **future** insured **exposure periods** (claim D).

# Three strategies for non-life reserving


## Aggregate reserving



## Granular reserving



## Individual reserving



# Three strategies for non-life reserving (continued)

## Aggregate reserving

First, we aggregate the past claim history into a **small number of summary statistics**, then we predict the total reserve based on these summary statistics.

## Granular reserving

First, we aggregate the past claim history into a **large number of summary statistics**, then we predict the total reserve based on these summary statistics.

## Individual reserving

First, we predict the future costs for **individual claims**, then we aggregate these individual predictions to predict the total reserve.

# Data set used in this session

We illustrate reserving analytics based on a **simulated dataset** registering the **detailed development** of 30,000 claims between January 1, 2010 and December 31, 2019. Data is available in a daily and yearly format.

**Daily** data:

- with claim events recorded at daily precision
- with records correspond to the reporting of a claim or a payment being made.

The data is stored in a .rds file in the `data` directory in the course material:

```
# install.packages("rstudioapi")
dir <- dirname(rstudioapi::getActiveDocumentContext())$path
setwd(dir)

reserving_daily <- readRDS("data/reserving_data_daily.rds")
```

##	accident_number	accident_date	reporting_date	settlement_date	reporting_delay	settlement_delay	payment_date	payment_size
## 1	1	2012-12-02	2012-12-07	2013-03-10	5	98	2012-12-07	0.0000
## 2	1	2012-12-02	2012-12-07	2013-03-10	5	98	2012-12-13	446.2251
## 3	1	2012-12-02	2012-12-07	2013-03-10	5	98	2012-12-24	490.1792
## 4	1	2012-12-02	2012-12-07	2013-03-10	5	98	2013-01-06	587.7409
## 5	1	2012-12-02	2012-12-07	2013-03-10	5	98	2013-01-24	1726.0902
## 6	2	2014-02-04	2014-02-10	2014-08-21	6	198	2014-02-10	0.0000



# Data set used in this session

We illustrate reserving analytics based on a **simulated dataset** registering the **detailed development** of 30,000 claims between January 1, 2010 and December 31, 2019. Data is available in a daily and yearly format.

**Yearly** data:

- aggregated into one record per claim per development year.

The data is stored in a .rds file in the `data` directory in the course material:

```
# install.packages("rstudioapi")
dir <- dirname(rstudioapi::getActiveDocumentContext())$path
setwd(dir)

reserving_yearly <- readRDS("data/reserving_data_yearly.rds")
```

```
## # A tibble: 6 x 10
##   accident_number accident_year reporting_year settlement_year development_year calendar_year size payment close is_closed
##         <int>         <dbl>         <dbl>         <dbl>         <dbl>         <dbl> <dbl> <dbl> <dbl> <dbl>
## 1             1             2             2             3             1             2  936.     1     0     0
## 2             1             2             2             3             2             3 2314.     1     1     1
## 3             2             4             4             4             1             4 3537.     1     1     1
## 4             3             6             6             6             1             6 1975.     1     1     1
## 5             4             9             10            10             1             9    0       0     0     0
## 6             4             9             10            10             2            10 3107.     1     1     1
```



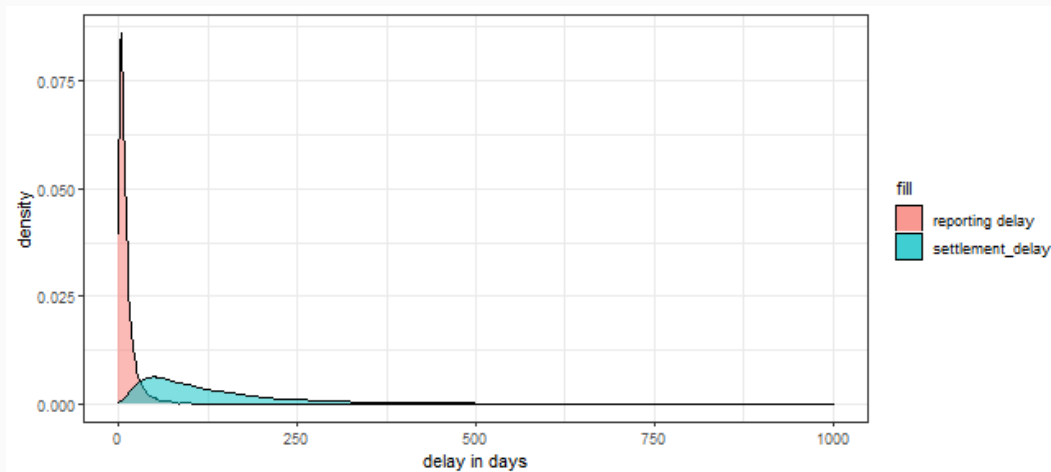
## Your turn

In this warm up exercise you load the `reserving_daily` or `reserving_yearly` data and get some feel for the data. Choose for each task the format that is most convenient (daily or yearly).

1. Visualize the reporting and settlement delay of claims with a density plot.
2. When was the last payment registered in the data set?  
What do you conclude?
3. What is the average number of payments per claim?
4. Calculate the number of claims per accident year.

For **Q.1** we plot the density of reporting and settlement delay:

```
ggplot(data = reserving_daily) +  
  theme_bw() +  
  geom_density(aes(reporting_delay, fill = 'reporting delay'),  
               alpha = .5) +  
  geom_density(aes(settlement_delay, fill = 'settlement_delay'),  
               alpha = .5) +  
  xlab('delay in days') +  
  xlim(c(0, 1000))
```



For **Q.2** the last payment was in 2024, i.e. the simulated data is not yet censored!

```
max(reserving_daily$payment_date)
```

```
## [1] "2025-05-03"
```

For **Q.3** we average the number of payments per claim. Count only records with `payment_size > 0` to exclude the reporting event.

```
reserving_daily %>%
  group_by(accident_number) %>%
  summarise(payments = sum(payment_size > 0)) %>%
  ungroup() %>%
  summarise(average = mean(payments))
```

```
## # A tibble: 1 x 1
##   average
##   <dbl>
## 1     2.38
```

For **Q.4** we group the data by accident year and summarise the number of claims.

```
library(tidyverse)
reserving_yearly %>%
  filter(development_year == 1) %>%
  group_by(accident_year) %>%
  summarise(num_claims = n())
```

```
## # A tibble: 11 x 2
##   accident_year num_claims
##   *           <dbl>     <int>
## 1             0       2743
## 2             1       2725
## 3             2       2810
## 4             3       2756
## 5             4       2705
## 6             5       2706
## 7             6       2703
## 8             7       2708
## 9             8       2640
## 10            9       2751
## 11           10       2391
```

We apply the filter `development_year == 1` to count each claim only once. Claim frequency is slightly lower in accident year 10.

# Censoring the data

The simulated data is not censored, i.e. we observe the full development of all claims that occur before the end of 2019. In practice the available data set will be censored and we will only observe:

- reported claims
- development information in past calendar years.

We censor the data assuming that we observe data until the end of 2020:

```
observed_daily <- reserving_daily %>%  
  filter(payment_date <= as.Date('2020-12-31'))  
  
unobserved_daily <- reserving_daily %>%  
  filter(payment_date > as.Date('2020-12-31'))  
  
observed_yearly <- reserving_yearly %>%  
  filter(calendar_year <= 10,  
         reporting_year <= 10)  
  
unobserved_yearly <- reserving_yearly %>%  
  filter(calendar_year > 10 | reporting_year > 10)
```

# IBNR and RBNS reserves

Reserving models predict the total payments in the unobserved part of the data:

```
reserve_actual <- sum(unobserved_yearly$size)
reserve_actual
```

```
## [1] 2149312
```

Some reserving methods split this reserve into

- IBNR: a reserve for Incurred, But Not (yet) Reported claims
- RBNS: a reserve for Reported, But Not (yet) Settled claims.

```
unobserved_yearly %>%
  mutate(reported = (reporting_year <= 10)) %>%
  group_by(reported) %>%
  summarise(reserve = sum(size))
```

```
## # A tibble: 2 x 2
##   reported reserve
## * <lgl>      <dbl>
## 1 FALSE     134124.
## 2 TRUE      2015188.
```

# Reserving data structures: individual and aggregated data

---

# From individual to aggregated data

## Individual triangle

The data set `reserving_yearly` describes the development of **individual claims discretised by calendar year**.

We can represent the claim characteristics in this data set (settlement, payment, payment size) with **two-dimensional tables** in which the rows represent **individual claims** and columns represent **development years**.

For payment size, this table consists of cells  $U_{k,j}^{\text{size}}$ , denoting the total amount paid for claim  $k$  in development year  $j$ .

## Runoff triangle

Many reserving models go one step further and remove individual claim characteristics by aggregating claims by occurrence year. The data is then represented in a (small) two-dimensional table, the so-called **incremental runoff triangle**.

$$X_{i,j}^{\text{size}} = \sum_{\text{occ.year}(\mathbf{k})=i} U_{k,j}^{\text{size}}.$$

As a result of the Law of Large Numbers (LLN), aggregating data from many claims into runoff triangles **reduces the variance** when all claims are independent realizations from the same distribution.



# From individual to aggregated data (continued)

We construct a runoff triangle by aggregating the individual claims by `accident_year`.

```
observed_yearly %>%  
  group_by(accident_year, development_year) %>%  
  summarise(value = sum(size)) %>%  
  pivot_wider(values_from = value,  
              names_from = development_year,  
              names_prefix = 'DY.')
```

`pivot_wider` creates for each unique value of `names_from` a new column

```
## # A tibble: 11 x 10  
## # Groups:   accident_year [11]  
##   accident_year    DY.1    DY.2    DY.3    DY.4    DY.5    DY.6    DY.7    DY.8    DY.9  
##   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1         0 4075901. 1610076. 112557. 24668. 10680. 8840. 9479. 6217. 3468.  
## 2         1 4037070. 1461939. 62622. 18518.    0    NA    NA    NA    NA  
## 3         2 4322579. 1774121. 105850. 41429. 24099. 7031.    0    NA    NA  
## 4         3 4162489. 1616297. 104487. 15982.    NA    NA    NA    NA    NA  
## 5         4 3903230. 1605598. 160906. 24194. 11622.    NA    NA    NA    NA  
## 6         5 4052230. 1623940. 131721. 26686. 14096. 6262.    NA    NA    NA  
## 7         6 4087125. 1593968. 96484. 40310. 4930.    NA    NA    NA    NA  
## 8         7 4059058. 1771141. 170375. 37317.    NA    NA    NA    NA    NA  
## 9         8 4033678. 1309634. 140624.    NA    NA    NA    NA    NA    NA  
## 10        9 4141278. 1615333.    NA    NA    NA    NA    NA    NA    NA  
## 11       10 3443223.    NA    NA    NA    NA    NA    NA    NA    NA
```

# From individual to aggregated data (continued)

A more sophisticated function for computing incremental runoff triangles is available in the R code.

```
incremental_triangle(observed_yearly,  
                    variable = 'payment',  
                    lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
##	[1,]	2054	523	17	6	2	1	1	1	1	0	0
##	[2,]	2074	482	16	4	0	0	0	0	0	0	NA
##	[3,]	2106	549	25	5	4	2	0	0	0	NA	NA
##	[4,]	2033	532	24	5	0	0	0	0	NA	NA	NA
##	[5,]	2008	473	32	7	2	0	0	NA	NA	NA	NA
##	[6,]	2035	498	26	4	2	2	NA	NA	NA	NA	NA
##	[7,]	2028	487	17	5	1	NA	NA	NA	NA	NA	NA
##	[8,]	2030	528	34	9	NA	NA	NA	NA	NA	NA	NA
##	[9,]	2005	444	24	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	2100	507	NA	NA	NA	NA	NA	NA	NA	NA	NA
##	[11,]	1784	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA

Using this function you can easily inspect the reserving data with multiple triangles.

```
incremental_triangle(observed_yearly,  
                    variable = 'size',  
                    lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	4075901	1610076	112556.58	24668.43	10679.546	8840.487	9478.77	6217.266	3468.364
##	[2,]	4037070	1461939	62622.23	18517.80	0.000	0.000	0.00	0.000	0.000
##	[3,]	4322579	1774121	105850.01	41429.20	24098.625	7030.788	0.00	0.000	0.000
##	[4,]	4162489	1616297	104486.62	15982.39	0.000	0.000	0.00	0.000	NA
##	[5,]	3903230	1605598	160905.57	24194.32	11622.131	0.000	0.00	NA	NA
##	[6,]	4052230	1623940	131721.42	26685.98	14095.805	6261.739	NA	NA	NA
##	[7,]	4087125	1593968	96483.79	40310.11	4930.009	NA	NA	NA	NA
##	[8,]	4059058	1771141	170374.72	37317.06	NA	NA	NA	NA	NA
##	[9,]	4033678	1309634	140624.40	NA	NA	NA	NA	NA	NA
##	[10,]	4141278	1615333	NA	NA	NA	NA	NA	NA	NA
##	[11,]	3443223	NA	NA	NA	NA	NA	NA	NA	NA

# Cumulative runoff triangles

**Cumulative runoff triangles** are constructed by computing the cumulative row sum of an **incremental runoff triangle**, i.e.

$$C_{i,j}^{\text{size}} = \sum_{l=1}^j X_{i,l}^{\text{size}}.$$

When using a new reserving method, check whether the incremental or cumulative runoff triangle is required!

```
incremental_triangle(observed_yearly,  
                    variable = 'size',  
                    lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	4075901	1610076	112556.58	24668.43	10679.546	8840.487	9478.77	6217.266	3468.364
##	[2,]	4037070	1461939	62622.23	18517.80	0.000	0.000	0.00	0.000	0.000
##	[3,]	4322579	1774121	105850.01	41429.20	24098.625	7030.788	0.00	0.000	0.000
##	[4,]	4162489	1616297	104486.62	15982.39	0.000	0.000	0.00	0.000	NA
##	[5,]	3903230	1605598	160905.57	24194.32	11622.131	0.000	0.00	NA	NA
##	[6,]	4052230	1623940	131721.42	26685.98	14095.805	6261.739	NA	NA	NA
##	[7,]	4087125	1593968	96483.79	40310.11	4930.009	NA	NA	NA	NA
##	[8,]	4059058	1771141	170374.72	37317.06	NA	NA	NA	NA	NA
##	[9,]	4033678	1309634	140624.40	NA	NA	NA	NA	NA	NA
##	[10,]	4141278	1615333	NA	NA	NA	NA	NA	NA	NA
##	[11,]	3443223	NA	NA	NA	NA	NA	NA	NA	NA

```
cumulative_triangle(observed_yearly,  
                   variable = 'size',  
                   lower_na = TRUE)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	4075901	5685977	5798533	5823202	5833881	5842722	5852200	5858418	5861886	5861886
##	[2,]	4037070	5499009	5561632	5580150	5580150	5580150	5580150	5580150	5580150	5580150
##	[3,]	4322579	6096700	6202550	6243979	6268078	6275109	6275109	6275109	6275109	NA
##	[4,]	4162489	5778787	5883273	5899256	5899256	5899256	5899256	5899256	NA	NA
##	[5,]	3903230	5508828	5669733	5693928	5705550	5705550	5705550	NA	NA	NA
##	[6,]	4052230	5676170	5807891	5834577	5848673	5854935	NA	NA	NA	NA
##	[7,]	4087125	5681093	5777577	5817887	5822817	NA	NA	NA	NA	NA
##	[8,]	4059058	5830199	6000574	6037891	NA	NA	NA	NA	NA	NA
##	[9,]	4033678	5343311	5483936	NA	NA	NA	NA	NA	NA	NA
##	[10,]	4141278	5756610	NA	NA	NA	NA	NA	NA	NA	NA
##	[11,]	3443223	NA	NA	NA	NA	NA	NA	NA	NA	NA

# Claims reserving with triangles

---

# Mack chain ladder method

Assumptions underneath Mack's (1993) **chain ladder method**:

1. There exist development factors  $f_j$  such that

$$E(C_{i,j+1} \mid C_{i,1}, \dots, C_{i,j}) = C_{i,j} \cdot f_j \quad \text{for } 1 \leq i \leq l, 1 \leq j \leq l-1,$$

with  $l$  the dimension of the runoff triangle.

2. There exist parameters  $\sigma_j$  such that

$$\text{Var}(C_{i,j+1} \mid C_{i,1}, \dots, C_{i,j}) = C_{i,j} \cdot \sigma_j^2 \quad \text{for } 1 \leq i \leq l, 1 \leq j \leq l-1.$$

3. Occurrence years are independent.

## Remarks:

- Method based on the **cumulative triangle**!
- Assumption 1. and 2. add a **Markov property**. Future evolutions of the reserve depend only on the last known situation!

# Mack chain ladder method (continued)

Estimating the **development factors**  $\hat{f}_j$ :

$$\hat{f}_j = \frac{\sum_{i=1}^{l-j} C_{i,j+1}}{\sum_{i=1}^{l-j} C_{i,j}}.$$

```
triangle <- cumulative_triangle(observed_yearly, variable = 'size')
l <- nrow(triangle)
f <- rep(0, l-1)

for(j in 1:(l-1)) {
  f[j] <- sum(triangle[1:(l-j), j+1]) / sum(triangle[1:(l-j), j])
}

f
```

```
## [1] 1.391002 1.021245 1.004906 1.001600 1.000630 1.000323 1.000263 1.000196 1.000000 1.000000
```

Development factors should converge to 1 for high development years.

# Mack chain ladder method (continued)

Use the development factors to estimate cells in the lower triangle (  $i + j > l + 1$  ):

$$\hat{C}_{i,j} = \hat{C}_{i,j-1} \cdot \hat{f}_{j-1} \quad \text{for } i + j > l + 1.$$

```
triangle_completed <- triangle
for(j in 2:l) {
  triangle_completed[l:(l-j+2), j] <- triangle_completed[l:(l-j+2), j-1] * f[j-1]
}
```

Completed cumulative triangle:

triangle\_completed

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	4075901	5685977	5798533	5823202	5833881	5842722	5852200	5858418	5861886	5861886
##	[2,]	4037070	5499009	5561632	5580150	5580150	5580150	5580150	5580150	5580150	5580150
##	[3,]	4322579	6096700	6202550	6243979	6268078	6275109	6275109	6275109	6275109	6275109
##	[4,]	4162489	5778787	5883273	5899256	5899256	5899256	5899256	5899256	5900411	5900411
##	[5,]	3903230	5508828	5669733	5693928	5705550	5705550	5705550	5707052	5708170	5708170
##	[6,]	4052230	5676170	5807891	5834577	5848673	5854935	5856829	5858371	5859518	5859518
##	[7,]	4087125	5681093	5777577	5817887	5822817	5826485	5828370	5829905	5831046	5831046
##	[8,]	4059058	5830199	6000574	6037891	6047551	6051361	6053318	6054913	6056098	6056098
##	[9,]	4033678	5343311	5483936	5510838	5519655	5523132	5524919	5526374	5527456	5527456
##	[10,]	4141278	5756610	5878910	5907750	5917202	5920930	5922845	5924405	5925565	5925565
##	[11,]	3443223	4789529	4891282	4915278	4923142	4926243	4927837	4929134	4930100	4930100

Completed incremental triangle:

```
require(ChainLadder)
cum2incr(triangle_completed)
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	4075901	1610076	112556.58	24668.43	10679.546	8840.487	9478.770	6217.266	3468.36	
##	[2,]	4037070	1461939	62622.23	18517.80	0.000	0.000	0.000	0.000	0.00	
##	[3,]	4322579	1774121	105850.01	41429.20	24098.625	7030.788	0.000	0.000	0.00	
##	[4,]	4162489	1616297	104486.62	15982.39	0.000	0.000	0.000	0.000	1155.08	
##	[5,]	3903230	1605598	160905.57	24194.32	11622.131	0.000	0.000	1502.662	1117.44	
##	[6,]	4052230	1623940	131721.42	26685.98	14095.805	6261.739	1893.935	1542.505	1147.07	
##	[7,]	4087125	1593968	96483.79	40310.11	4930.009	3667.976	1884.732	1535.009	1141.56	
##	[8,]	4059058	1771141	170374.72	37317.06	9660.235	3809.544	1957.475	1594.254	1185.56	
##	[9,]	4033678	1309634	140624.40	26902.60	8816.985	3477.005	1786.605	1455.090	1082.07	
##	[10,]	4141278	1615333	122299.67	28840.23	9452.018	3727.432	1915.283	1559.891	1160.06	
##	[11,]	3443223	1346306	101753.93	23995.21	7864.126	3101.242	1593.525	1297.837	965.13	

# Mack chain ladder method (continued)

The estimated reserve is the sum of the estimates in the **lower half** of the **incremental runoff triangle**.

```
triangle_completed_incr <- cum2incr(triangle_completed)
lower_triangle <- row(triangle_completed_incr) + col(triangle_completed_incr) > l+1

reserve_cl <- sum(triangle_completed_incr[lower_triangle])

data.frame(reserve_cl = reserve_cl,
            reserve_actual = reserve_actual,
            difference = reserve_cl - reserve_actual,
            relative_difference_pct = (reserve_cl - reserve_actual) / reserve_actual * 100)
```

```
##   reserve_cl reserve_actual difference relative_difference_pct
## 1    1734146      2149312  -415165.1          -19.31619
```



# Mack chain ladder method (continued)

These calculations are already implemented in the {ChainLadder} package.

```
require(ChainLadder)
triangle <- cumulative_triangle(observed_yearly, variable = 'size')
MackChainLadder(triangle)
```

```
## MackChainLadder(Triangle = triangle)
##
##      Latest Dev.To.Date Ultimate      IBNR Mack.S.E CV(IBNR)
## 1  5,861,886      1.000 5,861,886 0.00e+00      0      NaN
## 2  5,580,150      1.000 5,580,150 0.00e+00     635      Inf
## 3  6,275,109      1.000 6,275,109 9.31e-10     693 7.44e+11
## 4  5,899,256      1.000 5,900,411 1.16e+03    2,417 2.09e+00
## 5  5,705,550      1.000 5,708,170 2.62e+03    4,164 1.59e+00
## 6  5,854,935      0.999 5,859,518 4.58e+03    6,288 1.37e+00
## 7  5,822,817      0.999 5,831,046 8.23e+03    7,643 9.29e-01
## 8  6,037,891      0.997 6,056,098 1.82e+04   11,916 6.54e-01
## 9  5,483,936      0.992 5,527,456 4.35e+04   14,755 3.39e-01
## 10 5,756,610      0.971 5,925,565 1.69e+05   40,015 2.37e-01
## 11 3,443,223      0.698 4,930,100 1.49e+06  126,084 8.48e-02
##
##              Totals
## Latest:    61,721,361.87
## Dev:              0.97
## Ultimate: 63,455,508.27
## IBNR:         1,734,146.40
## Mack.S.E      136,785.03
## CV(IBNR):             0.08
```

**Latest**: Amount already paid

**ultimate**: Estimated total amount (= amount paid + reserve)

**IBNR**: Estimated total reserve, IBNR + RBNS!

**Mack.S.E**: Estimated standard deviation of the reserve

# GLM chain ladder method

The chain ladder reserve estimate can also be obtained by assuming a Poisson distribution for the incremental runoff triangle with multiplicative mean, i.e.

$$X_{i,j} \sim \text{POI}(\alpha_i \cdot \beta_j), \quad 1 \leq i, j \leq l.$$

```
triangle <- incremental_triangle(observed_yearly,  
                                variable = 'size')
```

```
triangle_long <- data.frame(  
  occ.year = as.numeric(row(triangle)),  
  dev.year = as.numeric(col(triangle)),  
  size = as.numeric(triangle))
```

```
head(triangle_long)
```

##	occ.year	dev.year	size
## 1	1	1	4075901
## 2	2	1	4037070
## 3	3	1	4322579
## 4	4	1	4162489
## 5	5	1	3903230
## 6	6	1	4052230

# GLM chain ladder method (continued)

The chain ladder reserve estimate can also be obtained by assuming a Poisson distribution for the incremental runoff triangle with multiplicative mean, i.e.

$$X_{i,j} \sim \text{POI}(\alpha_i \cdot \beta_j), \quad 1 \leq i, j \leq l.$$

This Poisson model can be estimated using the `glm` (Generalized Linear Model) routine in R:

$$\log(E(X_{i,j})) = \log(\alpha_i) + \log(\beta_j).$$

```
fit <- glm(size ~ factor(occ.year) + factor(dev.year),  
           data = triangle_long,  
           family = poisson(link = log))
```

`factor`: Treat the input as a categorical instead of a numeric variable.

```
summary(fit)
```

```
##  
## Call:  
## glm(formula = size ~ factor(occ.year) + factor(dev.year), family = poisson(link = log  
##      data = triangle_long)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max   
## -169.664   -55.719    -1.113    42.518   123.909   
##  
## Coefficients:  
##              Estimate Std. Error  z value Pr(>|z|)      
## (Intercept)   1.523e+01  4.220e-04 36079.449  <2e-16 ***  
## factor(occ.year)2 -4.926e-02  5.914e-04  -83.281  <2e-16 ***  
## factor(occ.year)3  6.812e-02  5.744e-04  118.589  <2e-16 ***  
## factor(occ.year)4  6.550e-03  5.832e-04   11.232  <2e-16 ***  
## factor(occ.year)5 -2.657e-02  5.881e-04  -45.183  <2e-16 ***  
## factor(occ.year)6 -4.040e-04  5.843e-04   -0.691    0.489   
## factor(occ.year)7 -5.275e-03  5.851e-04   -9.015  <2e-16 ***  
## factor(occ.year)8  3.259e-02  5.799e-04   56.205  <2e-16 ***  
## factor(occ.year)9 -5.874e-02  5.943e-04  -98.853  <2e-16 ***  
## factor(occ.year)10 1.080e-02  5.873e-04   18.397  <2e-16 ***  
## factor(occ.year)11 -1.731e-01  6.845e-04 -252.915  <2e-16 ***  
## factor(dev.year)2 -9.390e-01  2.950e-04 -3183.012  <2e-16 ***  
## factor(dev.year)3 -3.522e+00  9.734e-04 -3617.713  <2e-16 ***  
## factor(dev.year)4 -4.966e+00  2.096e-03 -2369.345  <2e-16 ***
```

# GLM chain ladder method

Predict the cells in the lower triangle

```
lower_triangle <- triangle_long$occ.year +  
  triangle_long$dev.year > 1 + 1  
  
triangle_long$size[lower_triangle] <-  
  predict(fit, newdata = triangle_long[lower_triangle, ], type = "size")  
  
triangle_long %>%  
  pivot_wider(values_from = size,  
              names_from = dev.year,  
              names_prefix = 'DY.')
```

```
## # A tibble: 11 x 12  
##   occ.year    DY.1    DY.2    DY.3    DY.4    DY.5    DY.6    DY.7    DY.8    DY.9    DY.10  
##   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1      1 4075901. 1610076. 112557. 24668. 10680. 8840. 9479. 6217. 3468. 0  
## 2      2 4037070. 1461939. 62622. 18518. 0 0 0 0 0 0  
## 3      3 4322579. 1774121. 105850. 41429. 24099. 7031. 0 0 0 0.000100  
## 4      4 4162489. 1616297. 104487. 15982. 0 0 0 0 1155. 0.0000941  
## 5      5 3903230. 1605598. 160906. 24194. 11622. 0 0 1503. 1117. 0.0000910  
## 6      6 4052230. 1623940. 131721. 26686. 14096. 6262. 1894. 1543. 1147. 0.0000934  
## 7      7 4087125. 1593968. 96484. 40310. 4930. 3668. 1885. 1535. 1142. 0.0000930  
## 8      8 4059058. 1771141. 170375. 37317. 9660. 3810. 1957. 1594. 1186. 0.0000965  
## 9      9 4033678. 1309634. 140624. 26903. 8817. 3477. 1787. 1455. 1082. 0.0000881  
## 10     10 4141278. 1615333. 122300. 28840. 9452. 3727. 1915. 1560. 1160. 0.0000945  
## 11     11 3443223. 1346306. 101754. 23995. 7864. 3101. 1594. 1298. 965. 0.0000786
```

Compute the reserve

```
reserve_glm <- sum(triangle_long$size[lower_triangle])  
reserve_glm
```

```
## [1] 1734146
```



## Your turn

We have demonstrated the estimation of the reserve (future amount paid). In this exercise you will use the same techniques to estimate the number of future payments.

1. Compute the actual number of future payments from the unobserved data set.
2. Create a cumulative triangle containing the number of payments per accident and development year.
3. Estimate the future number of payments with the chain ladder method implemented in the {ChainLadder} package.
4. Compute the difference between the estimated and actual number of payments. Express this error in terms of standard deviations.

For **Q.1** we compute the actual number of future payments from `unobserved_yearly`.

```
payment_actual <- sum(unobserved_yearly$payment)
payment_actual
```

```
## [1] 617
```

For **Q.2** we use `cumulative_triangle` with `variable = 'payment'` to create a payment triangle.

```
triangle <- cumulative_triangle(observed_yearly,
                                variable = 'payment')
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11]
## [1,] 2054 2577 2594 2600 2602 2603 2604 2605 2606 2606 2606
## [2,] 2074 2556 2572 2576 2576 2576 2576 2576 2576 2576 NA
## [3,] 2106 2655 2680 2685 2689 2691 2691 2691 2691 NA NA
## [4,] 2033 2565 2589 2594 2594 2594 2594 2594 NA NA NA
## [5,] 2008 2481 2513 2520 2522 2522 2522 NA NA NA NA
## [6,] 2035 2533 2559 2563 2565 2567 NA NA NA NA NA
## [7,] 2028 2515 2532 2537 2538 NA NA NA NA NA NA
## [8,] 2030 2558 2592 2601 NA NA NA NA NA NA NA
## [9,] 2005 2449 2473 NA NA NA NA NA NA NA NA
## [10,] 2100 2607 NA NA NA NA NA NA NA NA NA
## [11,] 1784 NA NA NA NA NA NA NA NA NA NA NA
```

For **Q.3** we apply the chain ladder method to this payment triangle.

```
require(ChainLadder)
cl <- MackChainLadder(triangle)
cl
```

```
## MackChainLadder(Triangle = triangle)
##
##      Latest Dev.To.Date Ultimate   IBNR Mack.S.E CV(IBNR)
## 1    2,606      1.000    2,606   0.000   0.000    NaN
## 2    2,576      1.000    2,576   0.000   0.121    Inf
## 3    2,691      1.000    2,691   0.000   0.125    Inf
## 4    2,594      1.000    2,594   0.330   0.677  2.0532
## 5    2,522      1.000    2,523   0.561   0.862  1.5354
## 6    2,567      1.000    2,568   0.769   0.998  1.2973
## 7    2,538      0.999    2,540   1.577   1.435  0.9101
## 8    2,601      0.999    2,604   3.200   2.069  0.6464
## 9    2,473      0.997    2,481   8.443   2.673  0.3166
## 10   2,607      0.987    2,640  33.472   7.535  0.2251
## 11   1,784      0.793    2,250 466.225  28.387  0.0609
##
##              Totals
## Latest:    27,559.00
## Dev:        0.98
## Ultimate: 28,073.58
## IBNR:       514.58
## Mack.S.E    30.08
## CV(IBNR):   0.06
```

For **Q.4** we compute the prediction error of the chain ladder method.

```
ultimate <- sum(cum2incr(cl$FullTriangle))
already_paid <- sum(cum2incr(cl$Triangle), na.rm = TRUE)

payment_cl <- ultimate - already_paid
sigma_cl <- as.numeric(cl$Total.Mack.S.E)

error = payment_actual - payment_cl

round(c(error = error,
        pct_error = error / payment_actual * 100,
        std.dev = error / sigma_cl),2)
```

##	error	pct_error	std.dev
##	102.42	16.60	3.41

For our data set, the chain ladder method significantly underestimates the number of future payments.

Actual triangle:

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
##	[1,]	2054	523	17	6	2	1	1	1	1	0	0
##	[2,]	2074	482	16	4	0	0	0	0	0	0	0
##	[3,]	2106	549	25	5	4	2	0	0	0	0	0
##	[4,]	2033	532	24	5	0	0	0	0	0	0	0
##	[5,]	2008	473	32	7	2	0	0	0	0	0	0
##	[6,]	2035	498	26	4	2	2	1	0	0	0	0
##	[7,]	2028	487	17	5	1	1	0	0	0	0	0
##	[8,]	2030	528	34	9	3	1	0	0	0	0	0
##	[9,]	2005	444	24	6	2	0	0	0	0	0	0
##	[10,]	2100	507	19	5	3	1	0	0	0	0	0
##	[11,]	1784	524	39	8	3	1	0	0	0	0	0

Predicted triangle:

##		dev										
##	origin	1	2	3	4	5	6	7	8	9		
##	1	2054	523.0	17.000000	6.0000000	2.0000000	1.00000000	1.00000000	1.00000000	1.00000000		
##	2	2074	482.0	16.000000	4.0000000	0.0000000	0.00000000	0.00000000	0.00000000	0.00000000		
##	3	2106	549.0	25.000000	5.0000000	4.0000000	2.00000000	0.00000000	0.00000000	0.00000000		
##	4	2033	532.0	24.000000	5.0000000	0.0000000	0.00000000	0.00000000	0.00000000	0.00000000	0.3295224	
##	5	2008	473.0	32.000000	7.0000000	2.0000000	0.00000000	0.00000000	0.00000000	0.2409938	0.3204066	
##	6	2035	498.0	26.000000	4.0000000	2.0000000	2.00000000	0.1976744	0.2453127	0.3261488		
##	7	2028	487.0	17.000000	5.0000000	1.0000000	0.8161821	0.1955041	0.2426194	0.3225679		
##	8	2030	528.0	34.000000	9.0000000	1.582905	0.8369510	0.2004790	0.2487932	0.3307761		
##	9	2005	444.0	24.000000	5.394067	1.508290	0.7974988	0.1910288	0.2370655	0.3151839		
##	10	2100	507.0	24.48796	5.739759	1.604952	0.8486084	0.2032713	0.2522584	0.3353832		
##	11	1784	437.7	20.86878	4.891454	1.367749	0.7231888	0.1732290	0.2149761	0.2858155		

# When the chain ladder method fails

---



# Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

We construct various triangles to assess this assumption:

Number of open claims at the start of the year:

```
triangle_open <- incremental_triangle(  
  observed_yearly %>%  
    mutate(open = calendar_year <= settlement_year),  
  variable = 'open')  
triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
##	[1,]	2743	722	28	6	2	1	1	1	1	0	0
##	[2,]	2725	698	24	4	1	0	0	0	0	0	NA
##	[3,]	2810	789	35	5	5	2	1	0	0	NA	NA
##	[4,]	2756	743	33	6	0	0	0	0	NA	NA	NA
##	[5,]	2705	677	46	9	3	0	0	NA	NA	NA	NA
##	[6,]	2706	689	34	8	2	2	NA	NA	NA	NA	NA
##	[7,]	2703	706	27	7	1	NA	NA	NA	NA	NA	NA
##	[8,]	2708	730	45	10	NA	NA	NA	NA	NA	NA	NA
##	[9,]	2640	647	27	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	2751	711	NA	NA	NA	NA	NA	NA	NA	NA	NA
##	[11,]	2328	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA

Number of open claims at the end of the year:

```
triangle_open_end <- incremental_triangle(  
  observed_yearly %>%  
    mutate(open_end = (calendar_year <= settlement_year) & (cl  
  variable = 'open_end')  
triangle_open_end
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
##	[1,]	722	28	6	2	1	1	1	1	0	0	0
##	[2,]	698	24	4	1	0	0	0	0	0	0	NA
##	[3,]	789	35	5	5	2	1	0	0	0	NA	NA
##	[4,]	743	33	6	0	0	0	0	0	NA	NA	NA
##	[5,]	677	46	9	3	0	0	0	NA	NA	NA	NA
##	[6,]	689	34	8	2	2	1	NA	NA	NA	NA	NA
##	[7,]	706	27	7	1	1	NA	NA	NA	NA	NA	NA
##	[8,]	730	45	10	4	NA	NA	NA	NA	NA	NA	NA
##	[9,]	647	27	7	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	711	30	NA	NA	NA	NA	NA	NA	NA	NA	NA
##	[11,]	644	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA

Less claims occurred during the last accident year, but of those that occurred many are still open.

# Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

We construct various triangles to assess this assumption:

Percentage of open claims for which a payment is made.

```
triangle_payment <- incremental_triangle(  
  observed_yearly,  
  variable = 'payment')
```

```
triangle_payment / triangle_open
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]
##	[1,]	0.7488152	0.7243767	0.6071429	1.0000000	1.0000000	1	1	1	1	NaN	NaN
##	[2,]	0.7611009	0.6905444	0.6666667	1.0000000	0.0000000	NaN	NaN	NaN	NaN	NaN	NaN
##	[3,]	0.7494662	0.6958175	0.7142857	1.0000000	0.8000000	1	0	NaN	NaN	NA	NA
##	[4,]	0.7376633	0.7160162	0.7272727	0.8333333	NaN	NaN	NaN	NaN	NA	NA	NA
##	[5,]	0.7423290	0.6986706	0.6956522	0.7777778	0.6666667	NaN	NaN	NA	NA	NA	NA
##	[6,]	0.7520325	0.7227866	0.7647059	0.5000000	1.0000000	1	NA	NA	NA	NA	NA
##	[7,]	0.7502775	0.6898017	0.6296296	0.7142857	1.0000000	NA	NA	NA	NA	NA	NA
##	[8,]	0.7496307	0.7232877	0.7555556	0.9000000	NA	NA	NA	NA	NA	NA	NA
##	[9,]	0.7594697	0.6862442	0.8888889	NA	NA	NA	NA	NA	NA	NA	NA
##	[10,]	0.7633588	0.7130802	NA	NA	NA	NA	NA	NA	NA	NA	NA
##	[11,]	0.7663230	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA

Average amount paid given a payment is made.

```
triangle_size <- incremental_triangle(  
  observed_yearly,  
  variable = 'size')
```

```
triangle_size / triangle_payment
```


##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
##	[1,]	1984.372	3078.539	6620.975	4111.405	5339.773	8840.487	9478.77	6217.266	3468.364
##	[2,]	1946.514	3033.069	3913.889	4629.451	NaN	NaN	NaN	NaN	NaN
##	[3,]	2052.507	3231.551	4234.001	8285.841	6024.656	3515.394	NaN	NaN	NaN
##	[4,]	2047.462	3038.153	4353.609	3196.478	NaN	NaN	NaN	NaN	NA
##	[5,]	1943.840	3394.499	5028.299	3456.332	5811.065	NaN	NaN	NA	NA
##	[6,]	1991.268	3260.923	5066.209	6671.495	7047.903	3130.870	NA	NA	NA
##	[7,]	2015.347	3273.035	5675.517	8062.022	4930.009	NA	NA	NA	NA
##	[8,]	1999.536	3354.434	5011.021	4146.339	NA	NA	NA	NA	NA
##	[9,]	2011.809	2949.626	5859.350	NA	NA	NA	NA	NA	NA
##	[10,]	1972.037	3186.061	NA	NA	NA	NA	NA	NA	NA
##	[11,]	1930.058	NA	NA	NA	NA	NA	NA	NA	NA

# Detecting when the chain ladder method fails

The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

We inspect the (granular) daily claims data:

```
claims <- observed_daily %>%  
  group_by(accident_number) %>%  
  slice(1) %>%  
  ungroup()  
  
occ_intensity <- claims %>%  
  group_by(accident_date) %>%  
  summarise(count = n())  
  
require(zoo)  
occ_intensity$moving_average <-  
  rollmean(occ_intensity$count, 30, na.pad = TRUE)  
  
ggplot(occ_intensity) +  
  theme_bw() +  
  geom_point(aes(x = accident_date, y = count)) +  
  geom_line(aes(x = accident_date, y = moving_average),  
            size = 1, color = 'blue') +  
  ggtitle('Evolution of claim frequency')
```



The moving average indicates a period with decreased claim frequency in 2020.

# Detecting when the chain ladder method fails


The chain ladder method assumes that claims from all occurrence years follow the **same development pattern**.

We inspect the (granular) daily claims data:

```
require(lubridate)
claims <- claims %>%
  mutate(start_year = floor_date(accident_date, unit = 'year'),
         time = as.numeric(accident_date - start_year) / 366,
         accident_year = year(accident_date),
         reporting_year = year(reporting_date)) %>%
  filter(accident_year == reporting_year)

ggplot(claims) +
  theme_bw() +
  geom_density(aes(x = time,
                  group = factor(accident_year),
                  color = factor(accident_year)))
```

Density showing claim occurrences within each accident year. The chain ladder method works best when the densities look similar across accident years.



The occurrence intensity in accident year 2020 deviates from the other accident years!

# Detecting when the chain ladder method fails

The data contains a drop in claim frequency during March-June 2020.

This emulates a situation where **less claims occurred** due to the first **COVID-19 lockdown**. Such changes in the claim occurrence and development pattern can severely impact chain ladder performance.

Other situations that might impact the performance of the chain ladder are:

- Inflation:
  - to be detected via the triangle `size`.
- Changing product mix:
  - to be detected via the triangles `close`, `payment` and `size`.
- Internal changes in claim handling:
  - to be detected via the triangles `close` and `payment`.
- Extreme weather events:
  - to be detected via the yearly distribution of occurrence dates.

# Fixing the chain ladder method

Improve the performance of the chain ladder method by making the **data more homogeneous**:

- **group** claims by:
  - claim characteristics, e.g. a separate triangle for fire and water related claims in home insurance.
  - occurrence time within a calendar year
- **omit**:
  - deviating years/cells, especially if they occurred long ago
- **adjust** for:
  - inflation.

The appropriate method(s) depend on the data inspection from the previous slides.

# Fixing the chain ladder method (continued)

Applied to the **COVID-19 reserving data set**:

- claim frequency has changed in the months March, April, May
- only 2020, i.e. accident year 10, is affected.

As 2020 is a recent calendar year, we cannot omit accident year 10 from the data.

We illustrate three approaches to improve the homogeneity across the accident years:

- applying the chain ladder method to more granular data
- grouping claims by occurrence month.


# Monthly chain ladder

A more **granular chain ladder** is more **robust** to **fluctuations in the occurrence process**.

```
require(lubridate)
claims <- observed_daily %>%
  group_by(accident_number) %>%
  slice(1) %>%
  ungroup() %>%
  mutate(start_month = floor_date(accident_date, unit = 'month',
    time = as.numeric(accident_date - start_month) / 31,
    accident_month = format(accident_date, '%Y%m'),
    reporting_month = format(reporting_date, '%Y%m')) %>%
  filter(accident_month == reporting_month)

ggplot(claims) +
  theme_bw() +
  geom_density(aes(x = time,
    group = factor(accident_month),
    color = factor(accident_month))) +
  theme(legend.position = 'none')
```

Within each month the claim occurrence density looks similar:



In real (non simulated) data, these monthly densities might be disturbed by seasonal effects, weekends and holidays.



# Monthly chain ladder

Construct a **monthly triangle** using the observed daily data:

```
triangle_month <- observed_daily %>%
  mutate(accident_month = year(accident_date)*12 + month(accident_date) - 2010*12,
         development_month = year(payment_date)*12 + month(payment_date) - 2010*12 - accident_month) %>%
  group_by(accident_month, development_month) %>%
  summarise(size = sum(payment_size)) %>%
  ungroup() %>%
  complete(expand.grid(accident_month = 1:132, development_month = 0:131), fill = list(size = 0)) %>%
  mutate(size = ifelse(accident_month + development_month > 132, NA, size)) %>%
  arrange(development_month) %>%
  pivot_wider(names_from = development_month, values_from = size) %>%
  arrange(accident_month)

triangle_month <- as.matrix(triangle_month[, 2:132])
```

# Monthly chain ladder

```
cl <- MackChainLadder(incr2cum(triangle_month))
summary(cl)$Totals
```

```
##              Totals
## Latest:      6.172136e+07
## Dev:         9.697826e-01
## Ultimate:    6.364453e+07
## IBNR:         1.923173e+06
## Mack S.E.:   1.035489e+05
## CV(IBNR):     5.384277e-02
```

Comparing the yearly and monthly chain ladder:

```
##           method cells reserve.estimate      error pct.error
## 1  Yearly chain ladder    66         1734146 -415165.1 -19.31619
## 2  Monthly chain ladder  8778         1923173 -226139.0 -10.52146
```

The monthly chain ladder predicts the reserve more accurately on this data set.

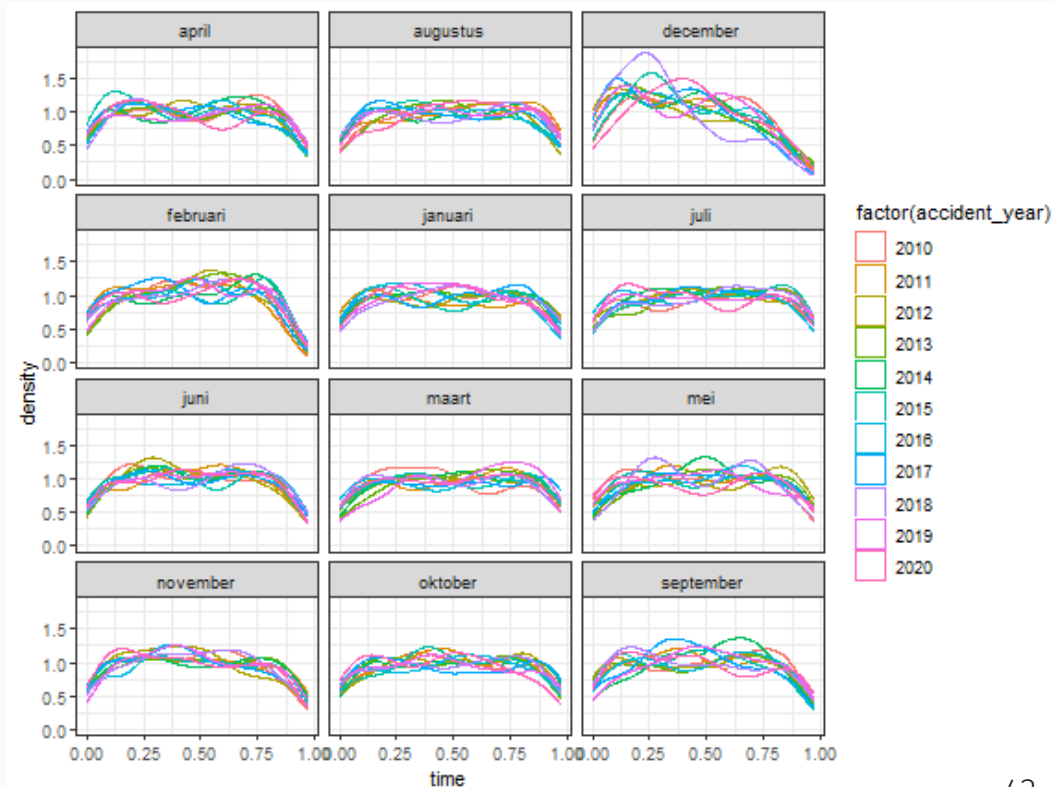
# Grouping claims by occurrence month

We group claims based on the occurrence month and estimate a **separate yearly chain ladder** for each of these groups.

```
require(lubridate)
claims <- observed_daily %>%
  group_by(accident_number) %>%
  slice(1) %>%
  ungroup() %>%
  mutate(start_month = floor_date(accident_date, unit = 'month',
    time = as.numeric(accident_date - start_month) / 31,
    accident_year = format(accident_date, '%Y'),
    reporting_year = format(reporting_date, '%Y'),
    month = format(accident_date, '%B')) %>%
  filter(accident_year == reporting_year)

ggplot(claims) +
  facet_wrap(~ month, ncol = 3) +
  theme_bw() +
  geom_density(aes(x = time,
    group = factor(accident_year),
    color = factor(accident_year)))
```

We are comparing occurrences in the same month accross different years. We are no longer comparing different months within one calendar year.



# Grouping claims by occurrence month

Add the accident date to the data set `reserving_yearly`

```
reserving_yearly <- reserving_yearly %>%  
  left_join(reserving_daily %>%  
    group_by(occident_number) %>%  
    slice(1) %>%  
    ungroup() %>%  
    select(occident_number, occident_date))  
  
reserving_yearly <- reserving_yearly %>%  
  mutate(occident_month = format(occident_date, '%B'))
```

Create a runoff triangle for each accident month

```
triangles <- reserving_yearly %>%  
  group_by(occident_month, occident_year, development_year) %>%  
  summarise(size = sum(size)) %>%  
  ungroup() %>%  
  complete(expand.grid(occident_month = unique(occident_month),  
    occident_year = 0:10, development_year =  
    fill = list(size = 0)) %>%  
  mutate(size = ifelse(occident_year + development_year > 11, 0, size))
```

```
triangles %>%  
  filter(occident_month == 'april') %>%  
  arrange(development_year) %>%  
  pivot_wider(names_from = development_year, values_from = size)
```

```
## # A tibble: 11 x 13  
##   occident_month occident_year  DY.1  DY.2  DY.3  DY.4  DY.5  DY.6  DY.7  DY.8  
##   <chr>          <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 april          0 406625. 70060. 8968. 4503. 2032.    0    0    0  
## 2 april          1 392370. 39237.    0    0    0    0    0    0  
## 3 april          2 490615. 62149.    0    0    0    0    0    0  
## 4 april          3 392533. 30396. 3977.    0    0    0    0    0  
## 5 april          4 426086. 40524. 19473. 5444. 5613.    0    0    NA  
## 6 april          5 423923. 43166. 3460.    0    0    0    NA    NA  
## 7 april          6 411784. 23560. 9287.    0    0    NA    NA    NA  
## 8 april          7 393665. 36667.    0    0    NA    NA    NA    NA  
## 9 april          8 394125. 30951.    0    NA    NA    NA    NA    NA  
## 10 april         9 388587. 24324.    NA    NA    NA    NA    NA    NA  
## 11 april        10 202582.    NA    NA    NA    NA    NA    NA    NA
```

# Grouping claims by occurrence month (continued)

We use the `glm` routine to apply the chain ladder method to each triangle:

`factor(development_year) : accident_month` estimates a different effect `factor(development)` for each `accident_month`

```
fit <- glm(size ~ factor(development_year) * accident_month + factor(accident_year) * accident_month,
          data = triangles,
          family = poisson(link = 'log'))

reserve_group <- sum(predict(fit, newdata = triangles %>% filter(is.na(size)), type = 'response'))
reserve_group
```

```
## [1] 1928050
```

We obtain a similar reserve estimate as with the monthly chain ladder


##	method	cells	reserve.estimate	error	pct.error
## 1	Yearly chain ladder	66	1734146	-415165.1	-19.31619
## 2	Monthly chain ladder	8778	1923173	-226139.0	-10.52146
## 3	Grouped chain ladder	792	1928050	-221261.9	-10.29455

# Research outlook


---

# Research outlook


## Aggregate reserving



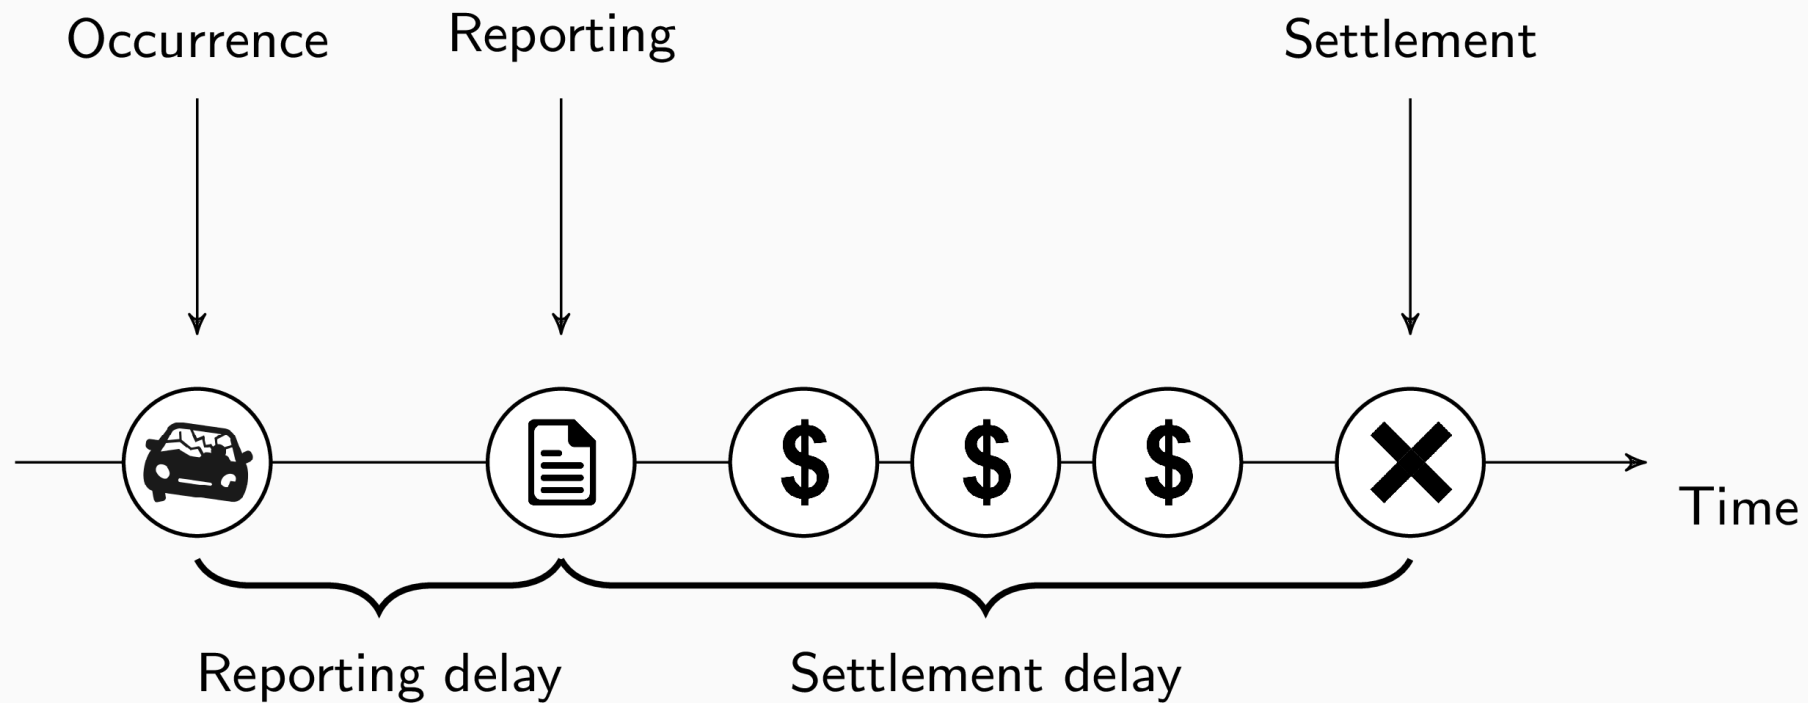
## Granular reserving



## Individual reserving

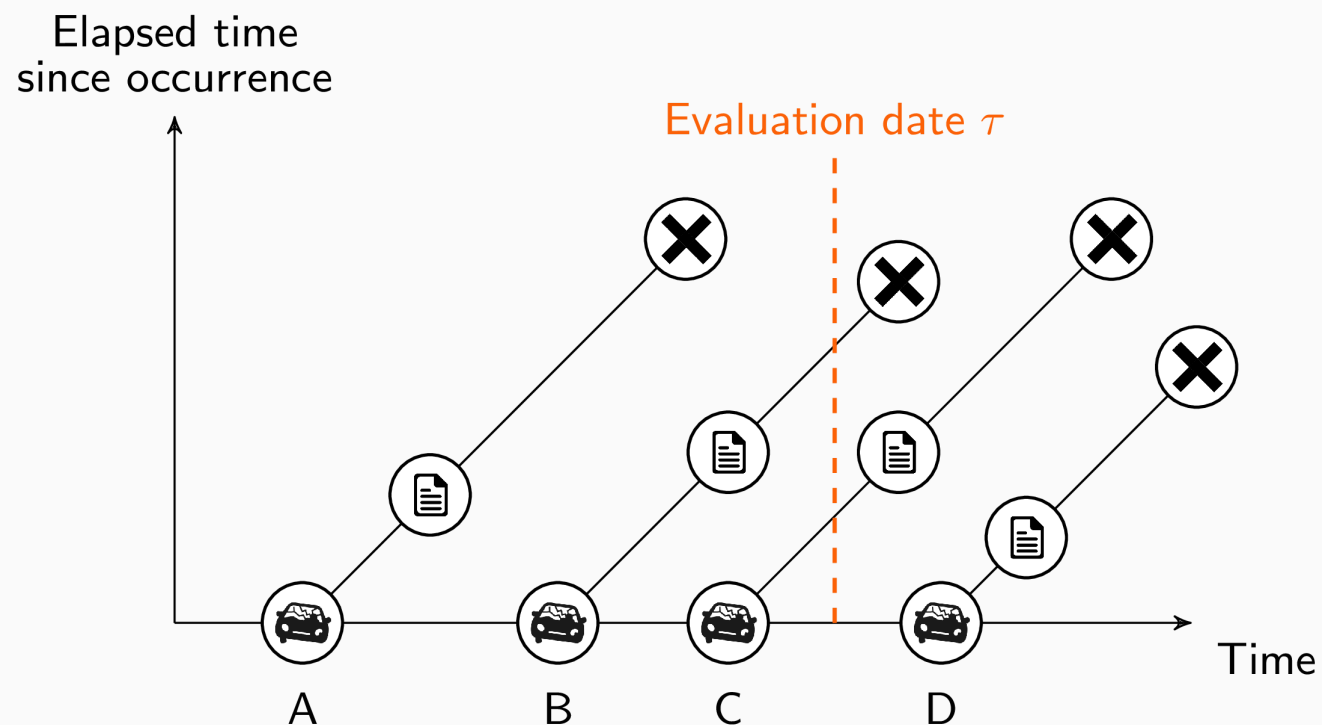


# Development of an individual claim





# Development of an individual claim (continued)



Observed claims are **censored** due to **delays** (reporting, settlement) in the claim development process.

**IBNR**: future costs for claims that occurred, but are not yet reported (claim B)

**RBNS**: future costs for claims that are reported, but are not yet settled (claim C)

**Pricing**: all costs for claims that will occur in **future** insured **exposure periods** (claim D).

# The IBNR reserve

Following a frequency-severity decomposition, the IBNR reserve is the product of the expected **number of unreported claims** times the **expected severity per claim**. Hereby:

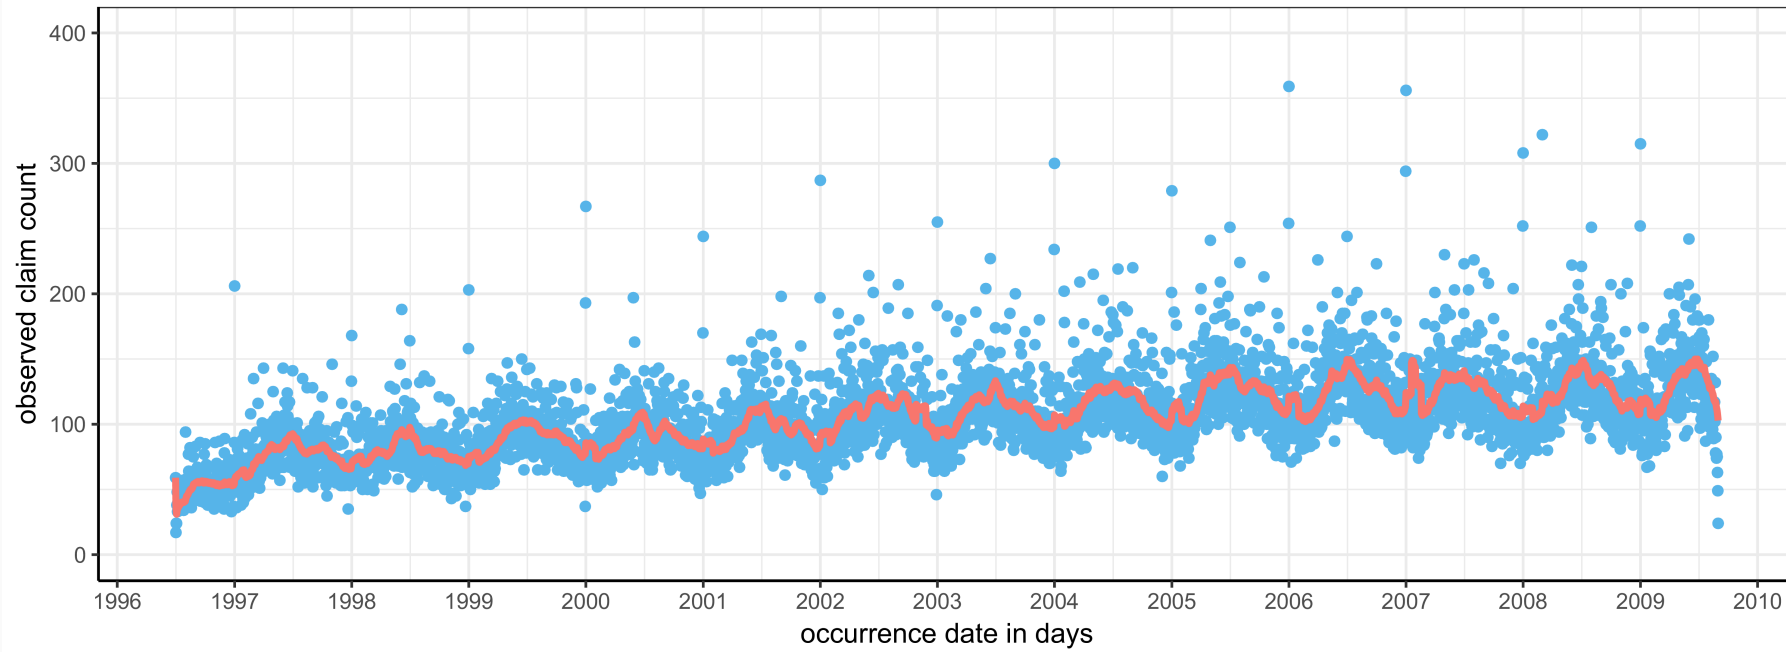
- insurance pricing covers the estimation of claim severity for new claims
- reserving methods focus on predicting the **number of unreported claims**.

Our work on the topic:

- "Modeling the number of hidden events subject to observation delay". Jonas Crevecoeur, Katrien Antonio and Roel Verbelen. (2019). European Journal of Operational Research.
- "Modeling the occurrence of events subject to a reporting delay via an EM algorithm". Roel Verbelen, Katrien Antonio, Gerda Claeskens and Jonas Crevecoeur. (2021). Statistical Science.

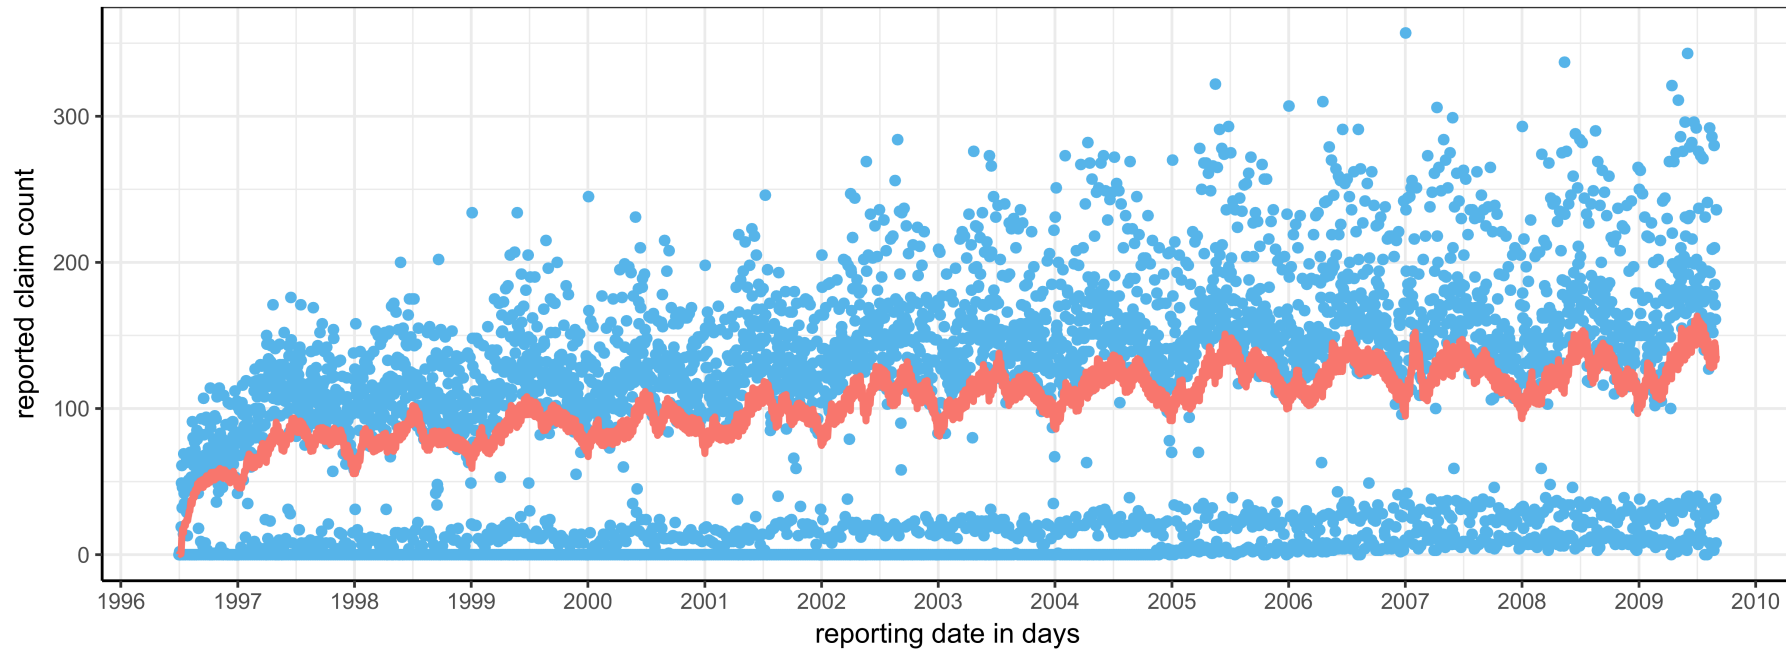
# The IBNR reserve (continued)

Number of observed claim occurrences per day



# The IBNR reserve (continued)

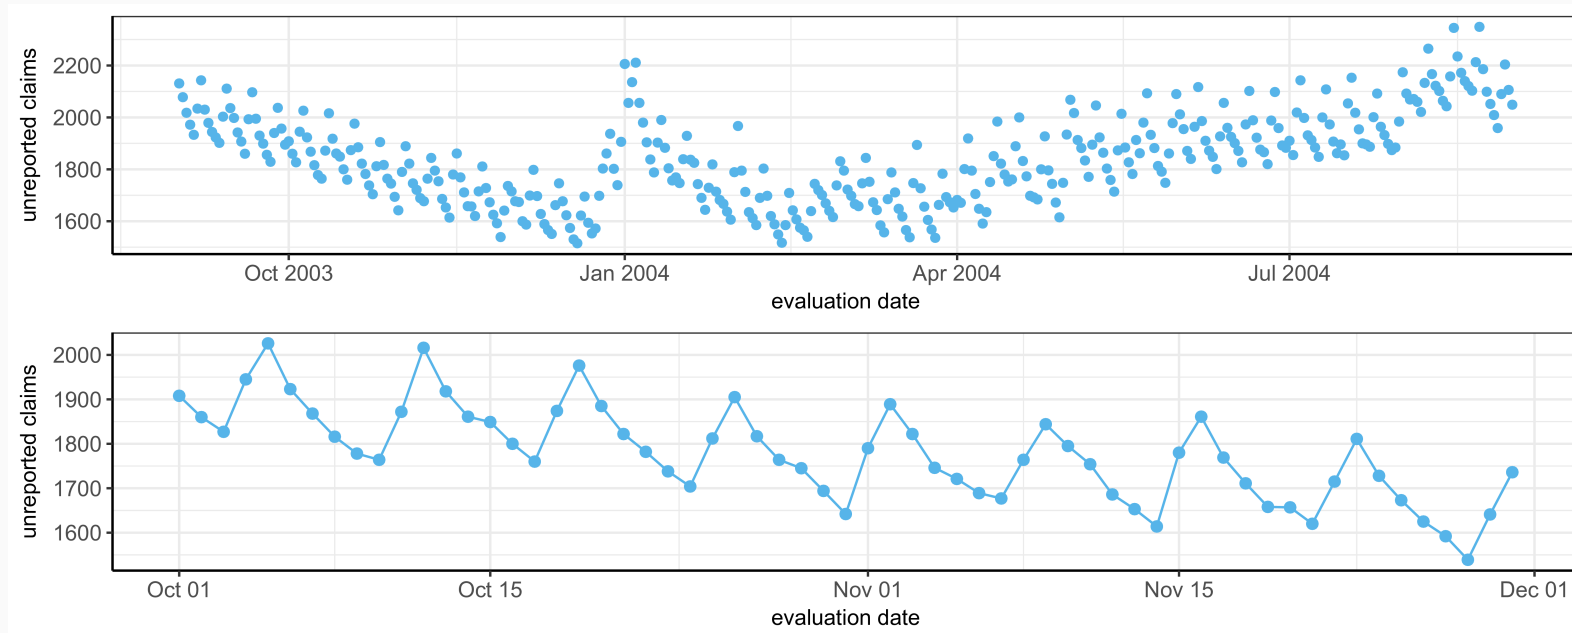
Number of reported claims per day



Almost no claims are reported on Saturdays, Sundays and holidays.

# The IBNR reserve (continued)

Total number of unreported claims on each day, i.e. the number of claims that occur before the evaluation date, but are reported afterwards.



The number of unreported claims increases during the weekend (+10% on sunday) and around the end of the year (+30%).

# The RBNS reserve

Predict the future costs of individual, reported claims.

Our work on the topic:

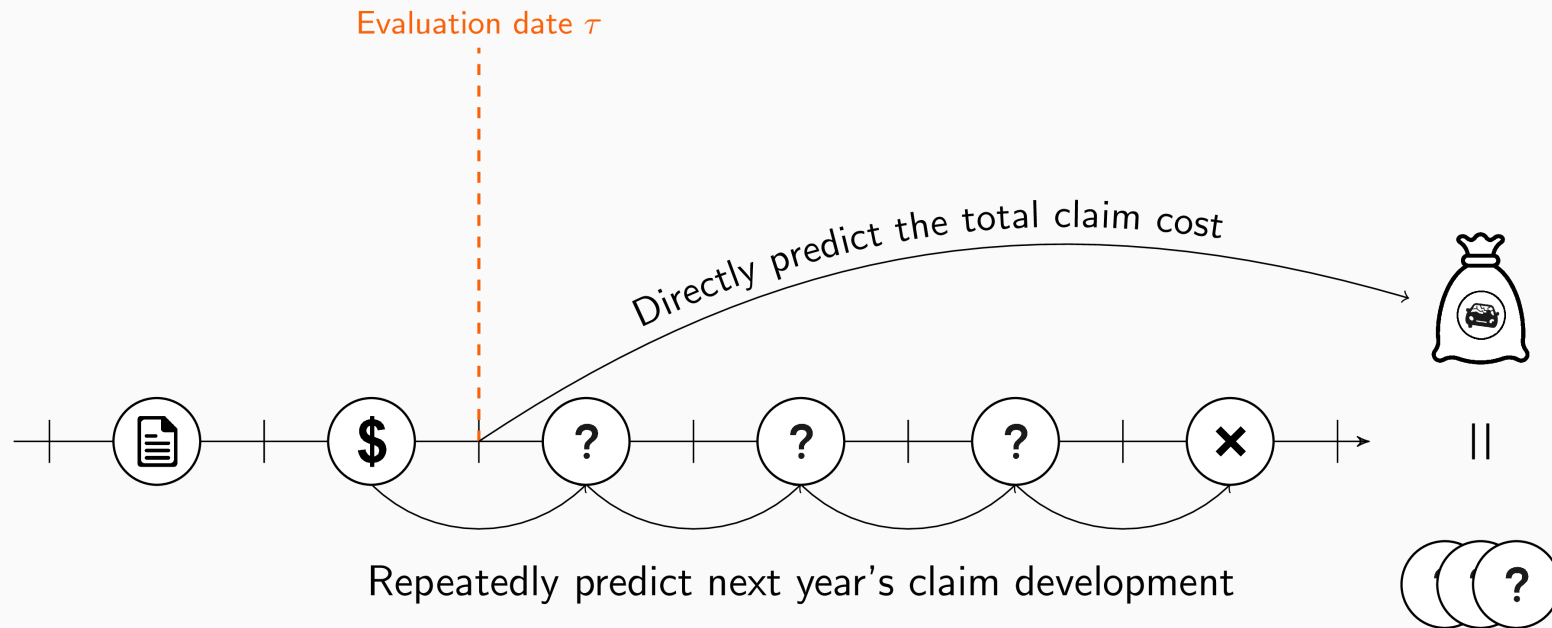
- "A hierarchical reserving model for non-life insurance claims". Jonas Crevecoeur and Katrien Antonio. (2021). Working paper, under review.
- "Bridging the gap between pricing and reserving with an occurrence and development model for non-life insurance claims". Katrien Antonio and Jonas Crevecoeur. (2021). Working paper.

Methods implemented in the R package {hirem}, covering tools for hierarchical reserving models, available on [GitHub](#).

# The RBNS reserve (continued)

The hierarchical model is based on **two key ideas**:

1. Model the development of individual claim in discrete time (steps of one year) using the observed development in previous years.



# The RBNS reserve (continued)

The hierarchical model is based on **two key ideas**:

1. Model the development of individual claim in discrete time (steps of one year) using the observed development in previous years.
2. Focus on all events registered over the lifetime of a claim

Common events registered over the lifetime of a claim:

- settlement
- payments
- initial incurred / changes in the incurred
- involvement lawyer.

These events are **dependent**:

- if a claim is **settled**, there will be no **payments** in the future
- large **payments** are more likely when the outstanding **reserve** is large.



# Implementation in the {hirem} package

Implementation readily available from the {hirem} package on [GitHub](#).

Events are added to the model as **layers**.

```
require(hirem)

individual_data <- individual_data %>%
  mutate(calendar_year = accident_year + development_year - 1,
         close = (settlement_year == calendar_year))

model <- hirem(individual_data %>% filter(calendar_year <= settlement_year)) %>%
  layer_glm('close', binomial(link = logit)) %>%
  layer_glm('payment', binomial(link = logit)) %>%
  layer_glm('size', Gamma(link = log),
         filter = function(data){data$payment == 1})

model <- fit(model,
  close = 'close ~ factor(development_year)',
  payment = 'payment ~ close + factor(development_year)',
  size = 'size ~ close + factor(development_year)')
```

# Thanks!

Slides created with the R package `xaringan`.

Course material available via

 <https://github.com/katrienantonio/workshop-loss-reserv-fraud>